

Introduction

Motivation – An OS is a piece of software that virtualises the computer; it provides a consistent and uniform view of underlying hardware and abstracts this complexity from the user, only the OS can access hardware directly.

Characteristics – (1) **Sharing**: share data, programs and hardware via time and space multiplexing. Resource allocation by ensuring fair and efficient use of resources, support simultaneous access to resources, provides mutual exclusion and protection against corruption. (2) **Concurrency**: support several simultaneous parallel activities, ability to switch activities at arbitrary times fairly, and ensure safe concurrency. (3) **Non-Determinism**: Never makes assumptions on events, as events arrive in unpredictable orders. (4) **Storing Data**: easy access to files through user defined names, enforce access controls, protection against failures, and storage management for easy expansion.

Structure – **Kernel**: core of the OS, implementing the most commonly executed functions of the OS, running in privileged mode (direct access to hardware), it is always in memory.

Monolithic Kernels – kernel is a single executable with its own address space, pushes parameters onto the stack and traps to execute system calls. **Advantages**: efficient calls within kernel and easy to write kernel components due to shared memory. **Disadvantages**: complex design and no protection between kernel components.


Microkernels – minimal kernel with majority of functionality in user space, engages in IPC with separate servers for I/O devices. **Advantages**: kernel not complex so less error prone, servers have clean interfaces and can crash and restart without bringing down the kernel. **Disadvantages**: high overhead of IPC communication.

Hybrid Kernels – Combines features of both monolithic and microkernels, it has a more structured design but offers a performance penalty for users.

Processes

Introduction – a process is the abstraction of running a program, it encapsulates the code and state of a program. It allows a single processor to run many processes simultaneously by virtualising the CPU. Processes can provide the illusion of concurrency, provide isolation of address space, a simplistic way of programming and allow better utilisation of machine resources.

Process Model – (1) **New**: process is being created. (2) **Ready**: process is runnable and waiting to be scheduled for the processor. (3) **Running**: process is currently executing on a processor. (4) **Waiting/Blocked**: process is waiting for an event to occur. (5) **Terminated**: process is being deleted.



Concurrency – Pseudo-concurrency is when a single hardware processor switches between processes over time by interleaving process execution, gives the illusion of concurrent execution. Real-concurrency is possible with multiple hardware processors, where each process is run on a separate CPU, still uses pseudo-concurrency as there are usually more processes than processors.

Multiprogramming – I/O is slow compared to computation time, so waiting for I/O is a poor utilisation of resources. Instead, multiprogramming allows the processor to be used by other programs while a program is performing I/O, allows the user to perform different activities at the same time.

Context Switches – switching between two processes due to an event (non-deterministic). When a context switch occurs, all process data is stored in a process control block, a data structure representing a process. PCB's are stored in a process table.

- On interrupt, the PC, PSW, some registers are pushed onto the (current) stack by the interrupt hardware
- Hardware jumps to address (PC from interrupt vector) to service interrupt
- Assembly language routine saves registers to PCB and then calls the device specific interrupt service routine
- C interrupt service runs (typically reads, writes & buffers data)
- Scheduler decides which process to run next
- C procedure returns control to assembly code
- Assembly procedure starts up the new current process

Cost – Context switches are expensive due to the direct cost of saving/restoring process state and the indirect cost of cache disruption and memory management registers. The information that must be stored are: (1) **Program counter**, page table register and other internal registers, (2) **process management** information such as process ID, priority and CPU uses, (3) **File management information** such as root directory, open directory, and open file descriptors.

Threads


Introduction – Threads are the smallest unit of execution that can be allocated processor time, they share the same address space as the parent process. Each thread has an execution state, saved thread context, execution stack, static storage and access to memory of the parent process.

Thread vs Processes – Processes are heavyweight: expensive to create and destroy, difficult to communicate between processes, blocking activities block entire application, and context switches are expensive. Threads are lightweight: faster creation and deletion, efficient communication between threads, activities can share data, and reflect parallelism within the application (blocking a thread does not block parent application). However threads can cause memory corruption and concurrency bugs.

User-Level Threads – Threads implemented by a software library in user space, kernel is oblivious to threads and thinks its managing the process only. The process itself handles thread scheduling and does not require kernel mode privileges. **Advantages**: better performance on creation and termination, switching, and synchronisation since it doesn't involve the kernel as it allows for application specific runtimes and scheduling algorithms. **Disadvantages**: blocking system call blocks all threads in the process (entire process blocks) since kernel is unaware of the existence of threads, defeats the whole purpose of threading. Page faults also block the entire process, and difficult to implement pre-emptive scheduling.

Kernel-Level Threads – Thread management is entirely done by the kernel. **Advantages**: Kernel can schedule multiple threads from the same process on different processors, page faults and system calls can be easily accommodated for and don't block the entire process. **Disadvantages**: thread creation/termination is more expensive as they involve kernel calls, but this can be mitigated using thread pools. Thread synchronisation is also expensive due to blocking system calls, same for scheduling.

Hybrid Approaches – Thread creation is done in user space which is multiplexed onto kernel threads, bulk of scheduling and synchronisation is done by the application.



Synchronisation

Critical Section – Section of code in which processes access a shared resource, this involves reads and writes at memory locations that are shared between processes/threads.

Race Condition – Occurs when multiple processes read and write data and the final result depends on the relative timing of their execution.

Solution – (1) **Mutual Exclusion**: no two processes access the critical section simultaneously, must first acquire permission. (2) **Progress**: no process running outside the critical section stops processes from entering the critical

section. (3) **Bounded waiting**: no process waits forever to access the critical section.

Disabling Interrupts – Using CLI() and STI() instructions to clear and set interrupt flags when a process enters the critical section. This only works on uniprocessor systems and does not guarantee bounded waiting.

Software Solution – Strict alternation between processes, does not guarantee progress. This is also called busy waiting, it wastes CPU time and should only be used when wait times are known to be short.

```
P0      turn 0      P1
while (true) {
  while (turn != 0)
    /* loop */ ;
  critical_section()
  turn = 1;
  noncritical_section0();
}


while (true) {
  while (turn != 1)
    /* loop */ ;
  critical_section()
  turn = 0;
  noncritical_section1();
}
```

Atomic Operations – Sequence of statements that appear to be indivisible. Examples of these are locks, which are set to 0 or 1 and prevent processes in accessing a critical section when it is busy. Lock granularity refers to the amount of data the lock is protecting. Locks that use busy waiting are called spin locks.

Lock Overhead – Measure of cost associated with using locks in terms of memory space, initialisation and time required to acquire and release locks. Lock contention is the measure of processes waiting for a lock (more contention means less parallelism). Coarser granularity means lower overhead and complexity but higher contention. Finer granularity means higher complexity and overhead but lower contention.

Semaphores – Processes cooperate by a means of signals: processes block waiting on a signal and continue if it has received a specific signal.

- Semaphores have two private components
- A counter (non-negative) integer
- A queue of processes currently waiting for that semaphore – queue is typically first in first out (FIFO)



Deadlocks

Conditions – A deadlock is when a set of processes are waiting for an event that only another process can cause.

- (1) **Mutual exclusion**: resource cannot be used by more than one process at a time.
- (2) **Hold and Wait**: process already holding resources may request a new resource.
- (3) **No pre-emption**: no resource can be forcibly removed from a process.
- (4) **Circular wait**: two or more processes in a circular chain, each waiting for a resource held by the next.

Strategies – Ignore it, detection and recovery, dynamic avoidance and prevention.

Detection and Recovery – Detect deadlocks and recover after occurrence. Build a directed resource ownership graph and perform and DFS to detect cycles. If deadlock detected then either temporarily remove resource from a process to break cycle, rollback to state previous to deadlock, or kill a process in the cycle.

Dynamic Avoidance – Bankers algorithm.

SAFE			UNSAFE		
	Has	Max		Has	Max
A	1	6	A	1	6
B	1	5	B	2	5
C	2	4	C	2	4
D	4	7	D	4	7
Free: 2			Free: 1		

Safe state

Are there enough resources to satisfy any (maximum) request from some customer? assume customer repays loan, and then check next customer closest to the limit, etc.

A state is safe iff there exists a sequence of allocations that guarantees that all customers can be satisfied

Prevention – (1) **Attack mutual exclusion**: share the resource. (2) **Attack hold and wait**: require all processes to request resources before start (violated non-determinism). (3) **Attack no pre-emption**: force process to give up resource half way through, not good in case of printer. (4) **Attack circular wait**: force single resource per process, if a process wants another resource it must release it first (optimality issues), difficult to organise large numbers of resources.

Livelock – Processes/threads are not blocked, but the system as a whole does not make progress.

Scheduling

Goals – Ensure fairness, avoid starvation, enforce policy, maximise resource utilisation, and minimise switching overhead.

Batch Systems – maximise throughput and CPU utilisation whilst minimising turnaround time.

Interactive Systems – Prioritise fast response times, and meets user expectations (predictability).


Real-time Systems – Meet hard and soft deadlines while being predictable.

Scheduling Types – Non-pre-emptive scheduling lets the process run until it blocks or voluntarily releases the CPU. Pre-emptive scheduling lets processes run for a fixed amount of time, utilise priority and clock interrupts.

CPU-Bound vs I/O-Bound Processes – CPU-bound processes spend most of their time using the CPU, whilst I/O bound processes spend most of their time using I/O and tend to use the CPU briefly in a short burst before issuing another I/O request.

FCFS Scheduling – Non-pre-emptive. **Advantages**: no indefinite postponement (all processes eventually scheduled), and easy to implement. **Disadvantages**: long jobs followed by short jobs drastically increase turnaround time (not fair), favours CPU bound processes over I/O and lower CPU utilisation.

Round-robin scheduling – Pre-emptive, process runs to completion or time quantum. **Advantages**: fair as all processes get same CPU time, good response times for small numbers of jobs. **Disadvantages**: low average turnaround time when run-times differ, and poor for similar run-times. Time quanta should be selected to be longer than context switch times, but provide a decent response time.



Shortest Job First – Non-pre-emptive with runtimes known in advance, optimal when all jobs available simultaneously.


Shortest Remaining Time – Pre-emptive version of SJF but run-times need to be known in advance, short jobs get good service, but may cause starvation of longer processes if short processes keep getting scheduled.

Priority Scheduling – Jobs run based on their priority, which can be user defined or based on process-specific metrics. Priorities can be static or dynamic.

General Purpose Scheduling – favours short I/O bound jobs, good resource utilisation and short response times. Quickly adapt to changes in process nature.

Multilevel Feedback Queues – implements aging, which increases a jobs priority as it waits to prevent starvation. They are not very flexible and do not react quickly to changes, cheating is also a concern and cannot donate priority.

- One queue for each priority level
- Run job on highest non-empty priority queue
- Each queue can use different scheduling algorithm
- Usually RR – quantum may be varied, e.g. highest priority is I/O-bound with short quantum. Exceed quantum → move down level but get bigger quantum



N-Step SCAN