

Summary: what we covered so far

- virtual memory has many uses
 - pretend we have lots of RAM
 - provide isolation between processes
 - provide sharing between processes
 - mapping of IO peripherals into CPU address space
- the page table maps virtual to physical pages
 - provides a fully associative mapping via table in RAM
- must translate virtual to physical on *every* access
 - all user-space addresses are virtual
 - even kernel-space addresses are virtual

TLB topics and common framework for memory hierarchy

(3rd Ed: p.527-555, 4th Ed: 502-534)

- TLB organisation: Intel and AMD processors
- TLB misses
- protection with virtual memory
- page faults
- common framework for memory hierarchy

Accelerating address translation

- page tables large: 10MB-1GB, so in main memory
 - problem: memory reference takes 2 memory accesses, to get physical address and to get data
- special cache for page table mappings:
 TLB (translation look-aside buffer)
- memory reference: look up virtual page number in TLB
 - if TLB hit: get physical page address
 - if TLB miss: load the relevant part of the page table into TLB and try again

TLB and page table

Virtual page number

TLB


Valid	Tag	Physical page address
1		
0		
1		
1		
1		
0		

Physical memory


Page table

Valid	Physical page or disk address
1	
1	
1	
1	
0	
1	
1	
0	
1	
1	
0	
1	

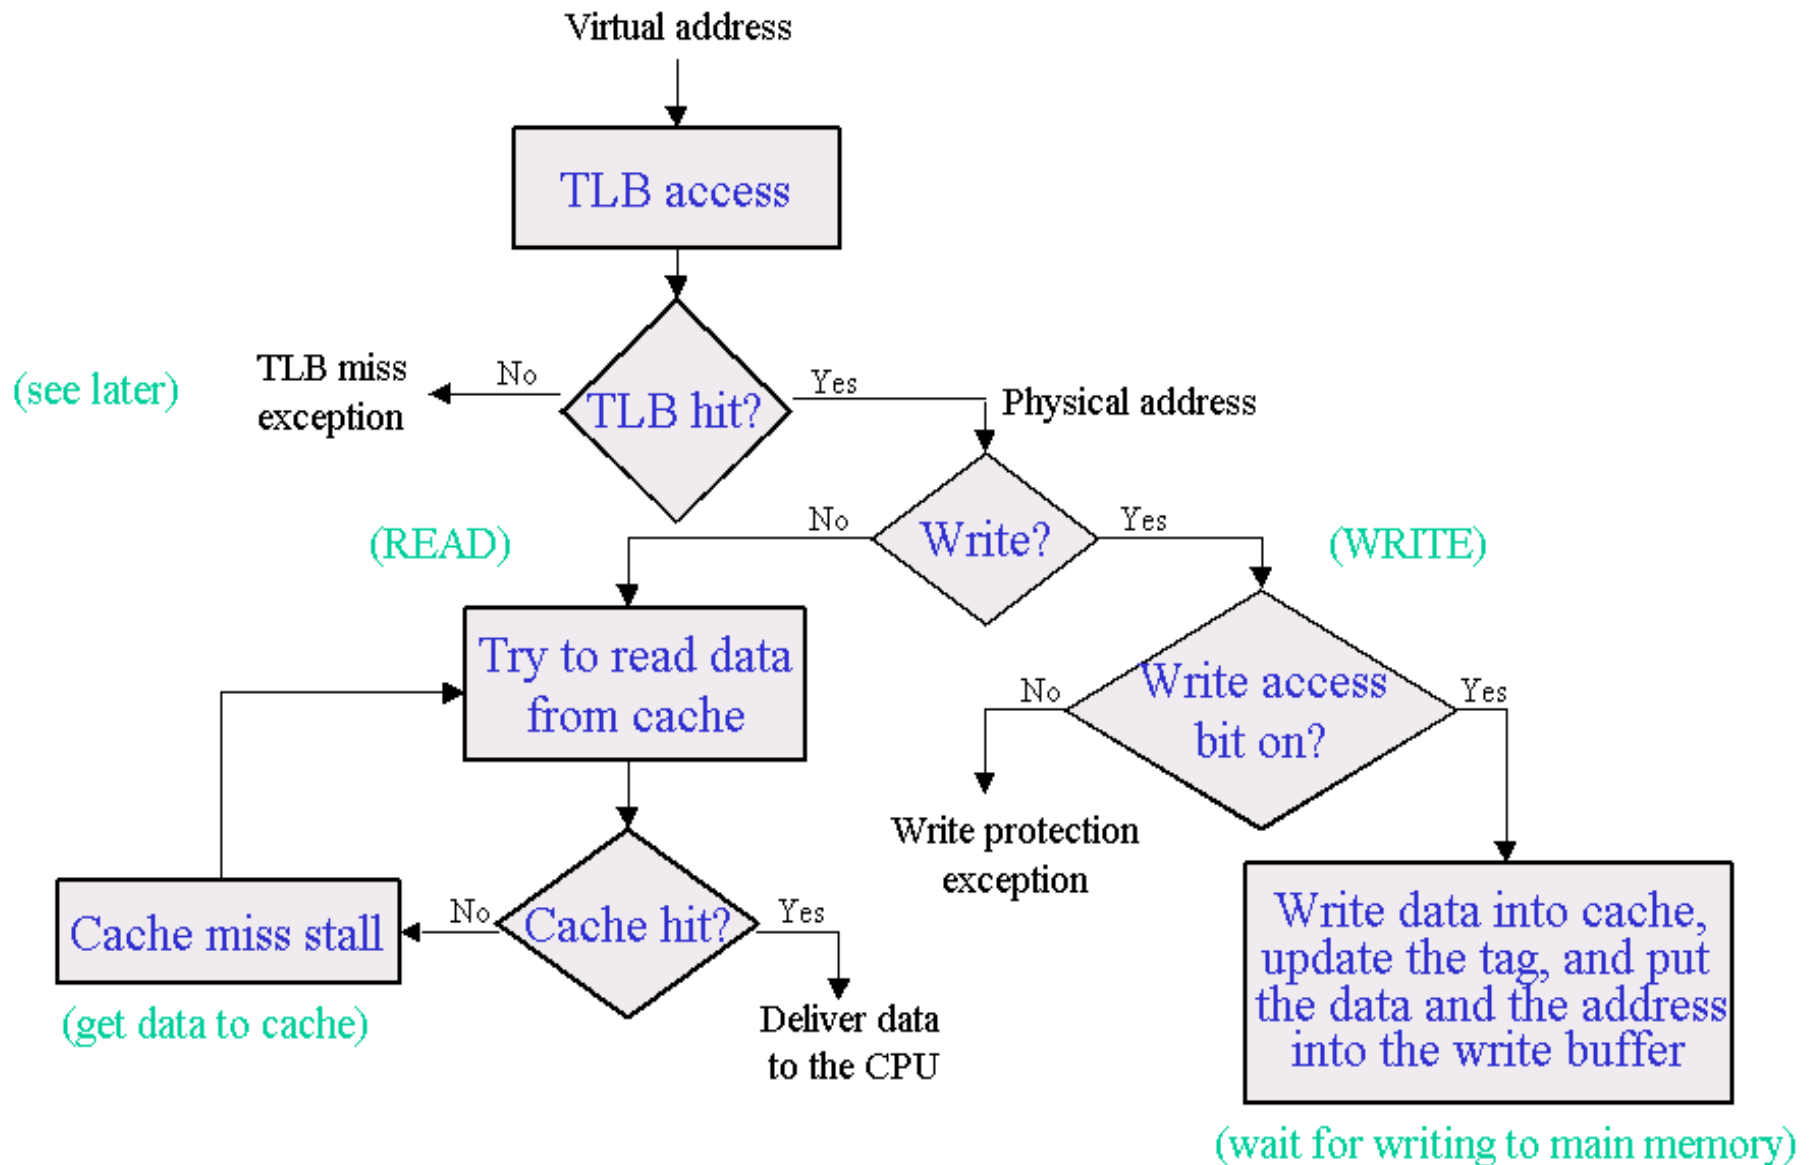
Disk storage



TLB and cache



Memory read and write



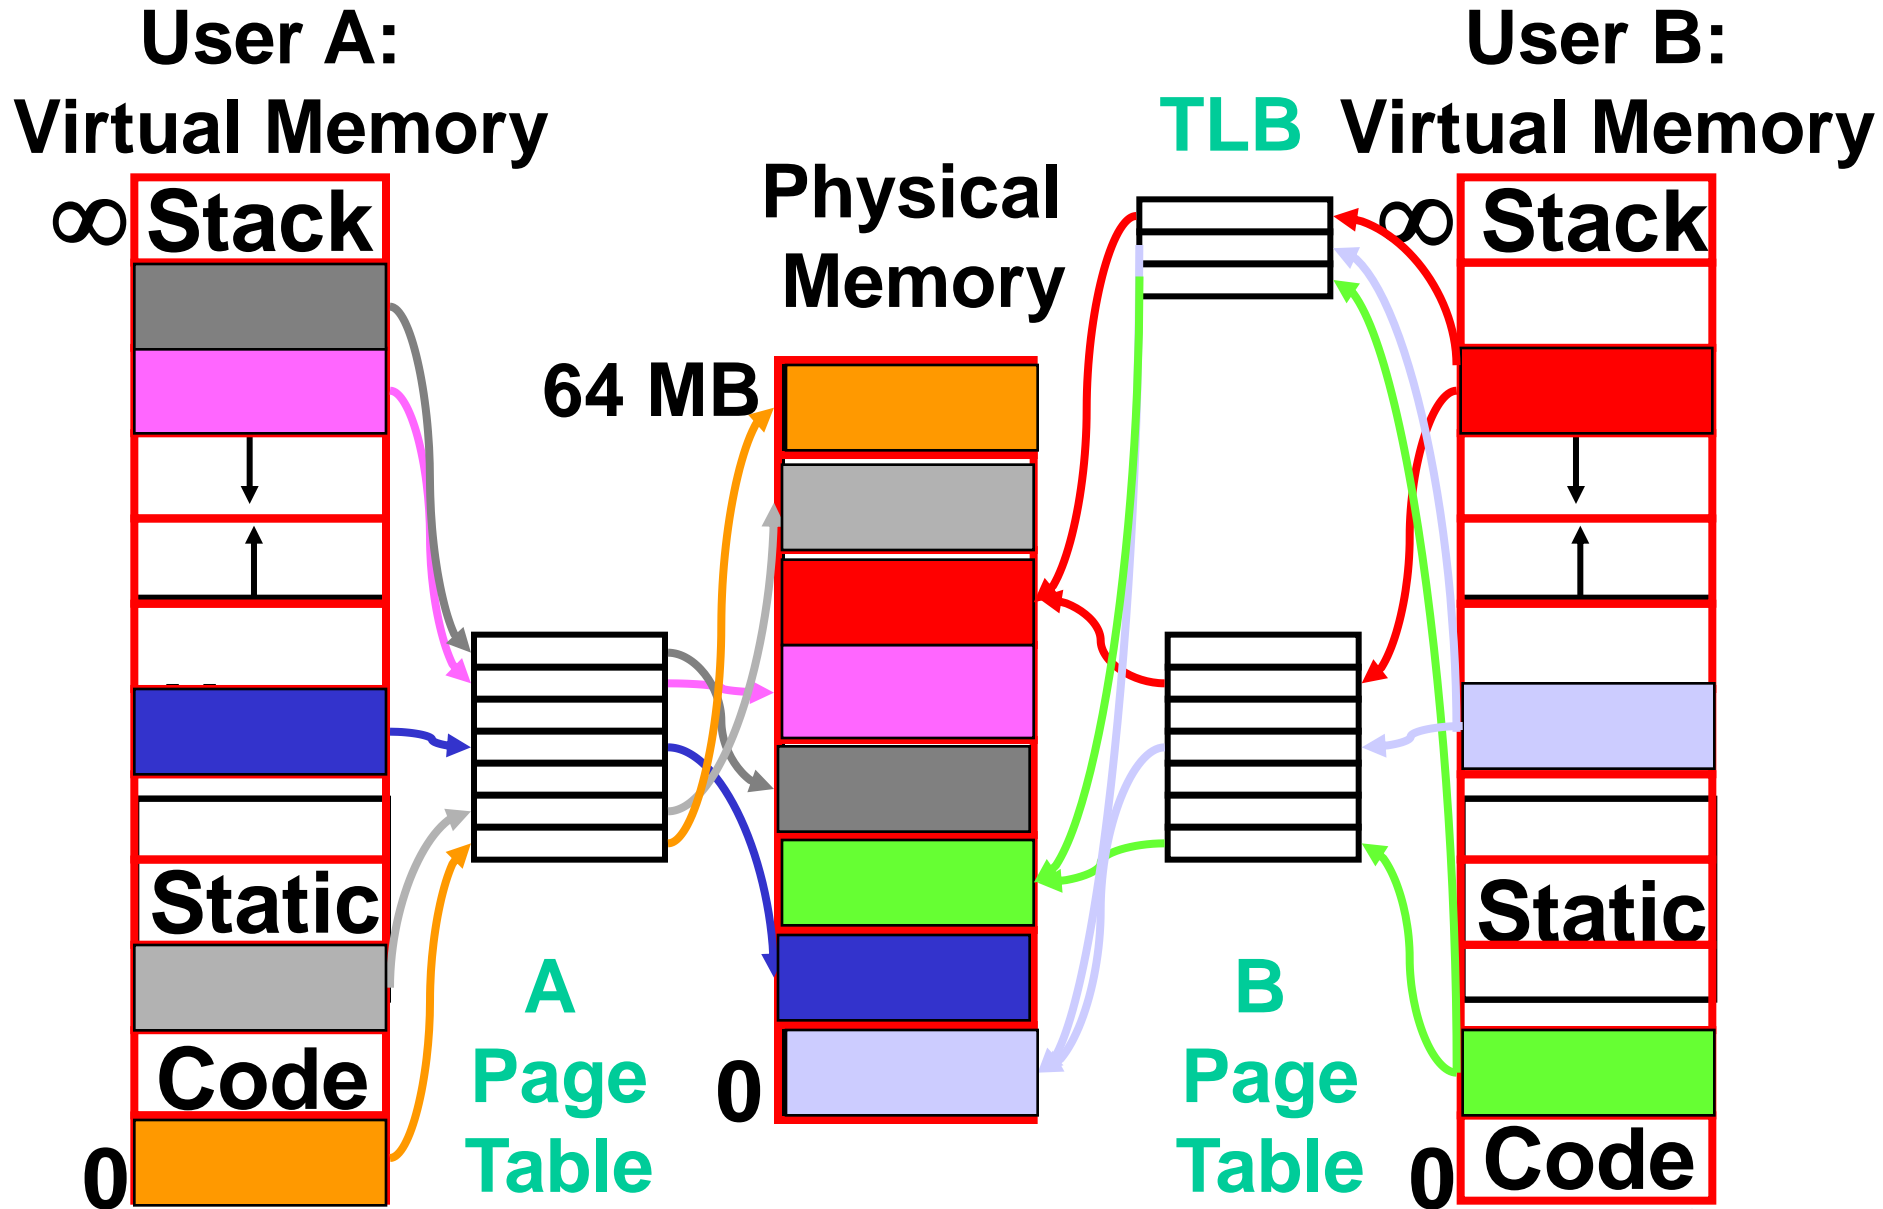
2-Level TLB Organization

	Intel Nehalem	AMD Opteron X4
Virtual addr	48 bits	48 bits
Physical addr	44 bits	48 bits
Page size	4KB, 2/4MB	4KB, 2/4MB
L1 TLB (per core)	L1 I-TLB: 128 entries for small pages, 7 per thread (2×) for large pages L1 D-TLB: 64 entries for small pages, 32 for large pages Both 4-way, LRU replacement	L1 I-TLB: 48 entries L1 D-TLB: 48 entries Both fully associative, LRU replacement
L2 TLB (per core)	Single L2 TLB: 512 entries 4-way, LRU replacement	L2 I-TLB: 512 entries L2 D-TLB: 512 entries Both 4-way, round-robin LRU
TLB misses	Handled in hardware	Handled in hardware

TLB miss

- no TLB entry matches virtual address
 - either page is in main memory,
just need to create missing TLB entry
 - or page is not in main memory,
page fault
- action: check valid bit of page table entry
 - if on, copy physical page address to TLB
 - if off, transfer control to operating system
(e.g. by exception) to deal with page fault
- quiz: what are possible combinations of hit/miss in cache, TLB and main memory?

Virtual memory: multiple processes



Protection with virtual memory

- a user process should not interfere with other users or the operating system
- page table can contain Access Right entries:
e.g. *Read Only*, *Read/Write*, *Executable*
in addition to the *valid* bit
- hardware supports 2 execution modes:
user mode and supervisor mode
- user process runs in users mode:
accesses its own page table by reading page table register (but not write page table register),
following access rights
- operating system runs in supervisor mode:
can change the page tables; change page table register and clear TLB for a new user process

Handling page faults

1. exception occurs, Cause Register indicates page fault
2. CPU changed to supervisor mode, disables further exceptions until enough saved state to recover
 1. EPC (Exception PC): address of offending instruction
 2. Cause Register: value indicates cause of exception
3. OS saves entire state of current process, e.g. EPC contains restart instruction when OS finishes dealing with page fault
 - avoid writing registers/memory - non-recoverable

Handling page faults (cont.)

4. find virtual address causing page fault:
 - in EPC, if instruction page fault
 - compute from base register and instruction offset, if data page fault (lw or sw)
5. once OS knows the virtual address:
 - look up page table to find disk address
 - choose physical page to replace; save replaced page on disk if dirty
 - copy referenced page from disk to memory
- OS finds another task for CPU during data transfer

Design parameters

Feature	Caches L1 / L2	paged memory	TLB
total size (blocks)	250-2000 / 15000-50000	16000-250000	40-1024
total size (kilobytes)	16-64 / 500-4000	10^6 - 10^9	0.25-16
block size (bytes)	16-64 / 64-128	4000-64000	4-32
miss penalty (cycles)	10-25 / 100-1000	10^7 - 10^8	10-1000
miss rate (global for L2)	2-5% / 0.1-2%	0.00001-0.0001%	0.01-2%

Common framework for memory hierarchy

- planning a block - upper level
one place, a few places, any place
- finding a block - upper level
by index, limited search, full search
- replacing a block - by one from lower level
random, LRU
- writing a block - to lower level
write through, write back

Placement schemes

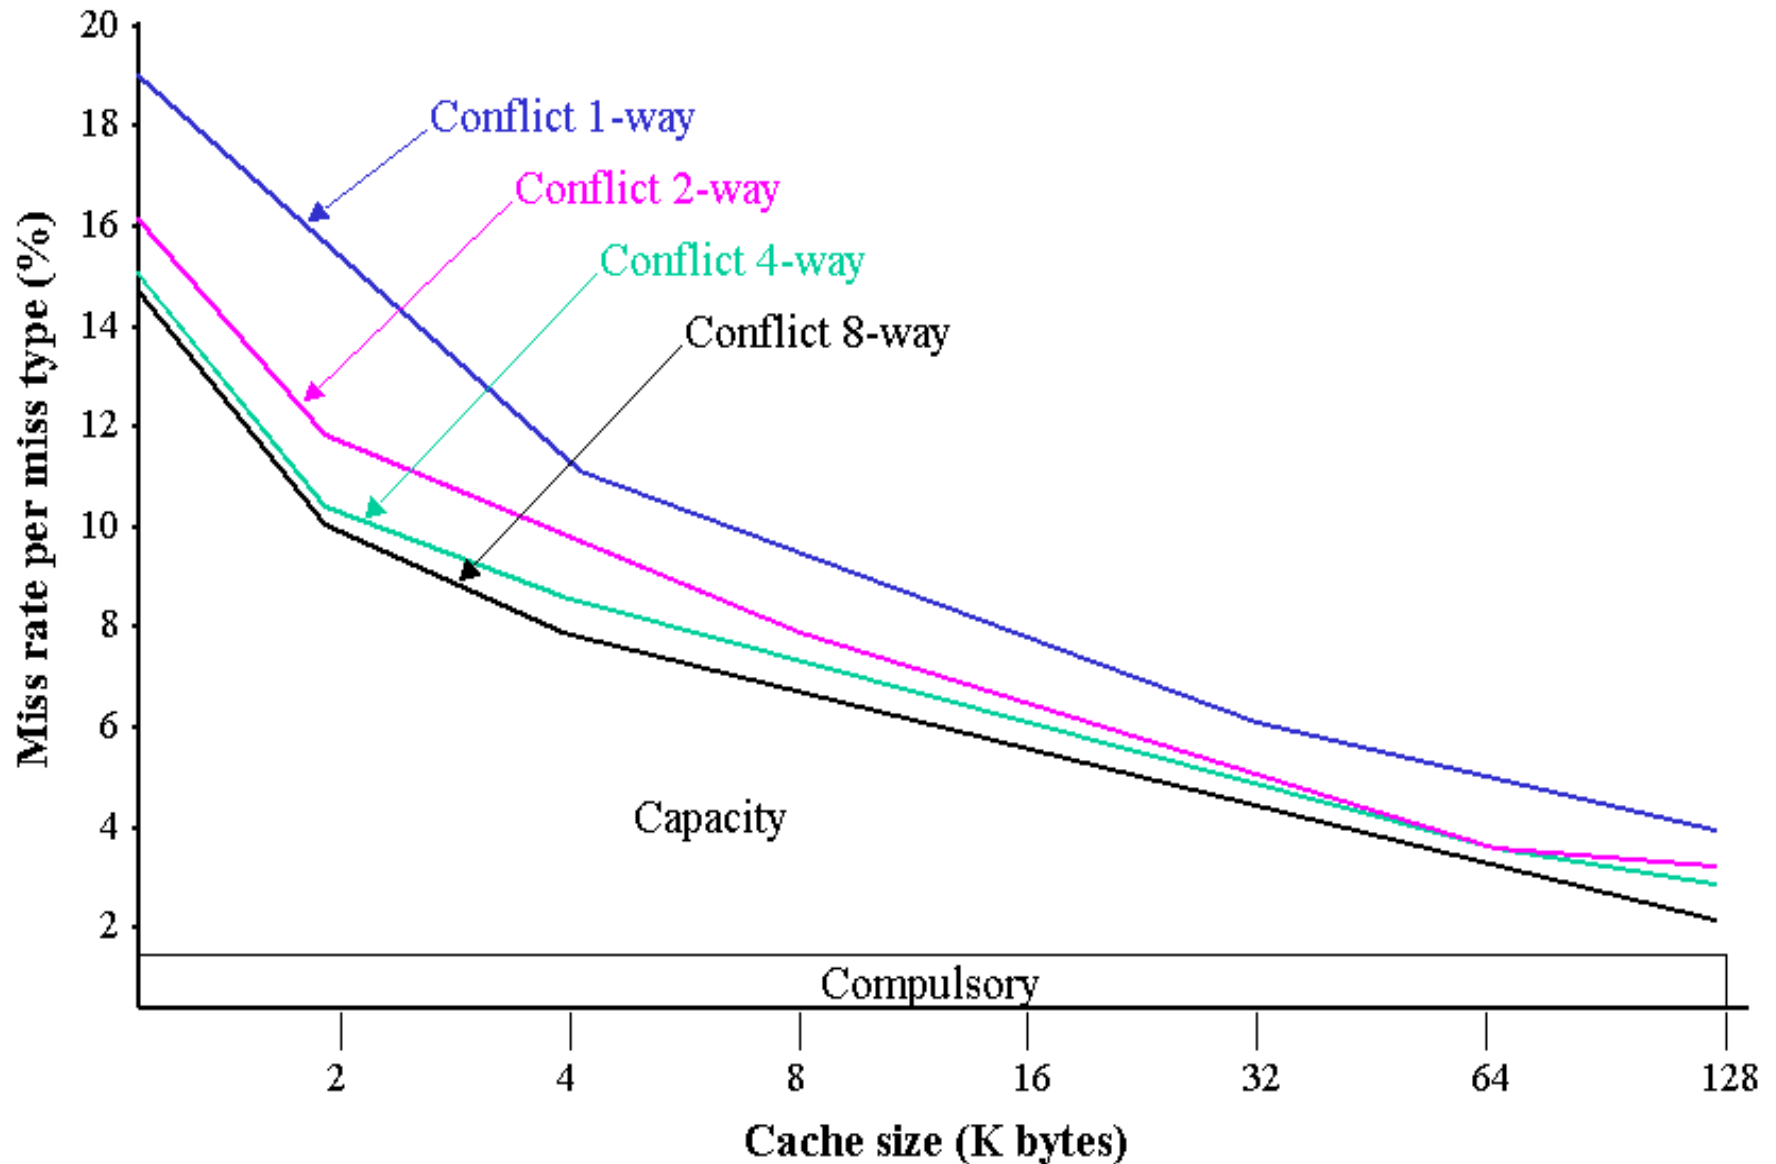
- *unique place* - direct mapped, e.g. cache
simple, fast access, high miss rate
- *any place* - fully associative, e.g. paged memory
low miss rate, costly, slow (search everywhere)
- *any n places* - n -way set associative, e.g. TLB
 - contains several sets, each n blocks
 - set with given block = (block number) **mod** (number of sets)
 - for fixed number of blocks, increased associativity leads to more blocks per set

Classify misses

- compulsory: first access to block after initialisation
- capacity: insufficient cache space
- conflict: rigid placement scheme

decision	effect on miss rate	possible negative effect
↑ size	↓ capacity misses	↑ access time, ↑ hardware
↑ associativity	↓ conflict misses	↑ access time, ↑ hardware
↑ block size	↓ compulsory misses	↑ miss penalty

Miss rate classification



Virtual memory: summary

- motivations: caching; running large programs; running multiple programs; protection
- page table: virtual to physical / disk address
- large page size: reduce miss rate
- place page anywhere: fully associative
- page replacement: LRU, write back, dirty bit
- TLB: page table cache


Summary: what to do on a write hit?


- write through
 - write to corresponding blocks at both upper and lower level
 - read misses do not require writes to the lower level
 - easy to implement
 - may require write buffer in high-speed systems
- write back
 - only write to upper level: at upper level speed
 - modify lower level block during replacement: minimise access overhead
 - better use of wide lower level, since write entire block
 - preferred choice when speed difference between the 2 levels is large, e.g. main memory and disk
 - more complicated to implement

Effect of cache size

- r : miss rate, d : data reference per instruction
 p : miss penalty (time) t : clock cycle time
 n : instruction count c : CPI without stall
- CPU time $\tau = (\text{CPU execution cycles} + \text{memory stall cycles}) \times t$

$$= (nc + n \times (r + rd) \times \lceil p/t \rceil) \times t$$


 instruction miss rate


 data miss rate
- 2 caches: large r_1, t_1 ; small r_2, t_2 $r_1 < r_2, t_1 > t_2$
- large cache slower if: $ct_1 + r_1p(1+d) > ct_2 + r_2p(1+d)$

$$c > \frac{p(1+d)(r_2-r_1)}{t_1-t_2}$$

Summary

	TLB	L1 cache	virtual memory
hit time	1 cycle	1-15 cycles	10-100 cycles
miss penalty	10-1000 cycles	10-25 cycles	10^7 - 10^8 cycles
placement	fully associative or set associative	direct mapped or set associative	fully associative
identification	tag / block	tag / block	table
replacement	random	direct mapped / random	LRU
write	flush on write to page table	write back / write through	write back