

Compilers I - Chapter 7:

Loop optimisations

- Lecturers:
 - Part I: Paul Kelly (phjk@doc.ic.ac.uk)
 - Office: room 304 William Penney Building
 - Part II: Naranker Dulay (nd@doc.ic.ac.uk)
 - Office: room 562
- Materials:
 - Textbook
 - Course web pages (<http://www.doc.ic.ac.uk/~phjk/Compilers>)
 - Piazza (<http://piazza.com/imperial.ac.uk/fall2016/221>)

The plan

- To optimise or not to optimise?
 - High-level vs low-level; role of analysis
 - Peephole optimisation
 - Local, global, interprocedural
 - Loop optimisations
 - Where optimisation fits in the compiler
 - Example: live ranges
 - Live ranges as a data flow problem
 - Solving the data-flow equations
 - Deriving the interference graph
 - Loop-invariant code and code motion optimisations
- Other data-flow analyses
 - More sophisticated optimisations

Loop-invariant code motion

- Definition:
 - An instruction is loop-invariant if its operands can only arrive from outside the loop
- Objective:
 - move (“hoist”) loop-invariant instructions out of loop
- Issues:
 - Where should we move the loop-invariant instructions *to*?
 - How can we find out whether operands only arrive from outside loop
 - Other pitfalls...

Finding loop-invariant instructions

- A CFG node is a definition if it updates a temporary
- In our CFG, an instruction can update at most one temporary, t
- Each definition node is labelled with the Node id, d :

$$d : t := u1 \oplus u2$$

Or simply

$$d : t := u1 \quad \text{or} \quad d : t := \text{constant}$$

(where $u1$ and $u2$ are given by the Node's "uses" field)

- This definition is loop-invariant if, for each $u_i \in \text{uses}(d)$,
 - All the definitions of u_i that reach d are outside the loop
 - Or only one definition of u_i reaches d , and that definition is loop invariant


Finding reaching definitions

- A definition of variable t is a statement which may assign to t
- A definition d **reaches** a program point p if there exists a path from d to p such that d is not killed along that path
- Consider a CFG node
$$n: \quad t := u1 \oplus u2 \quad (\text{defs}(n)=\{t\}, \text{uses}(n) = \{u1, u2\})$$

Define:

- **Gen(n)** is the set of definitions generated by node n , i.e. $\{n\}$
- **Kill(n)** is the set of all definitions of t , excluding n
- **ReachIn(n)** is the set of definitions reaching the point before n
- **ReachOut(n)** is the set of definitions reaching the point after n

Reaching definitions



- Reaching definitions link each use of a variable back to where its value was generated
- Loops and conditionals lead to multiple reaching definitions
- ((In the worst case, the number of reaching definitions could be quite large))


- $\text{Gen}(4) = \{4\}$
- $\text{kill}(4) = \{1\}$
- $\text{Gen}(5) = \{5\}$
- $\text{kill}(5) = \{\}$
- $\text{Gen}(6) = \{\}$

Reaching definitions – another data flow analysis

- Dataflow equations:

$$\mathbf{ReachIn}(n) = \bigcup_{p \in \text{pred}(n)} \mathbf{ReachOut}(p)$$

$$\mathbf{ReachOut}(n) = \mathbf{Gen}(n) \cup (\mathbf{ReachIn}(n) - \mathbf{Kill}(n))$$



(“*The $\mathbf{Gen}(n)$ + whatever survives* ”)


- Many dataflow problems have “gen” and “kill”
- In the case of $\mathbf{ReachOut}(n)$, $\text{gen}(n)$ is usually just its own id, $\{n\}$
- But if node n defines no value (eg it’s a jump), it will never reach anything – so $\text{gen}(n) = \{ \}$

Reaching definitions – another data flow analysis

- Dataflow equations:

$$\mathbf{ReachIn}(n) = \bigcup_{p \in \text{pred}(n)} \mathbf{ReachOut}(p)$$

$$\mathbf{ReachOut}(n) = \mathbf{Gen}(n) \cup (\mathbf{ReachIn}(n) - \mathbf{Kill}(n))$$



- Solve in the usual way: (“*The $\mathbf{Gen}(n)$ + whatever survives*”)
 - Initialise **ReachIn**(n) and **ReachOut**(n) to $\{ \}$
 - Iterate, repeatedly updating **ReachIn**(n) and **ReachOut**(n) using definitions above
 - Until convergence
 - At each step, the sets increase in size

Use reaching definitions to find loop invariant instructions


- Find the set of definitions of variables used by this node
- An instruction is loop invariant if the definitions of all the values it uses are outside the loop

- Example:

```
1  x = 1
2  w=100
3  z=200
```

Here:

```
4  x = x+1
5  y=w+z
6  if (x<10) goto Here
```




- Reaching definitions (ReachIn):

- 1: []
- 2: [1]
- 3: [1,2]
- 4: [1,2,3,4,5]
- 5: [2,3,4,5]
- 6: [2,3,4,5]

Use reaching definitions to find loop invariant instructions

- Find the definitions which reach this node which are relevant
 - that is, which generate the values this node uses:



- Reaching definitions:

- 1: []
- 2: [1]
- 3: [1,2]
- 4: [1,2,3,4,5]
- 5: [2,3,4,5]
- 6: [2,3,4,5]

- “Relevant” Reaching definitions:

- 1: []
- 2: []
- 3: []
- 4: [1,4]
- 5: [2,3]
- 6: [4]


All the definitions of the values used by node 5 lie outside the loop

Where should we move the loop-invariant instructions *to*?

- Given control-flow graph, need to find
 - Where the loops are
 - Where the loop headers are
 - So we can find a place to put the loop's loop-invariant instructions
 - Need robust scheme that handles all loops including goto
- Definition:

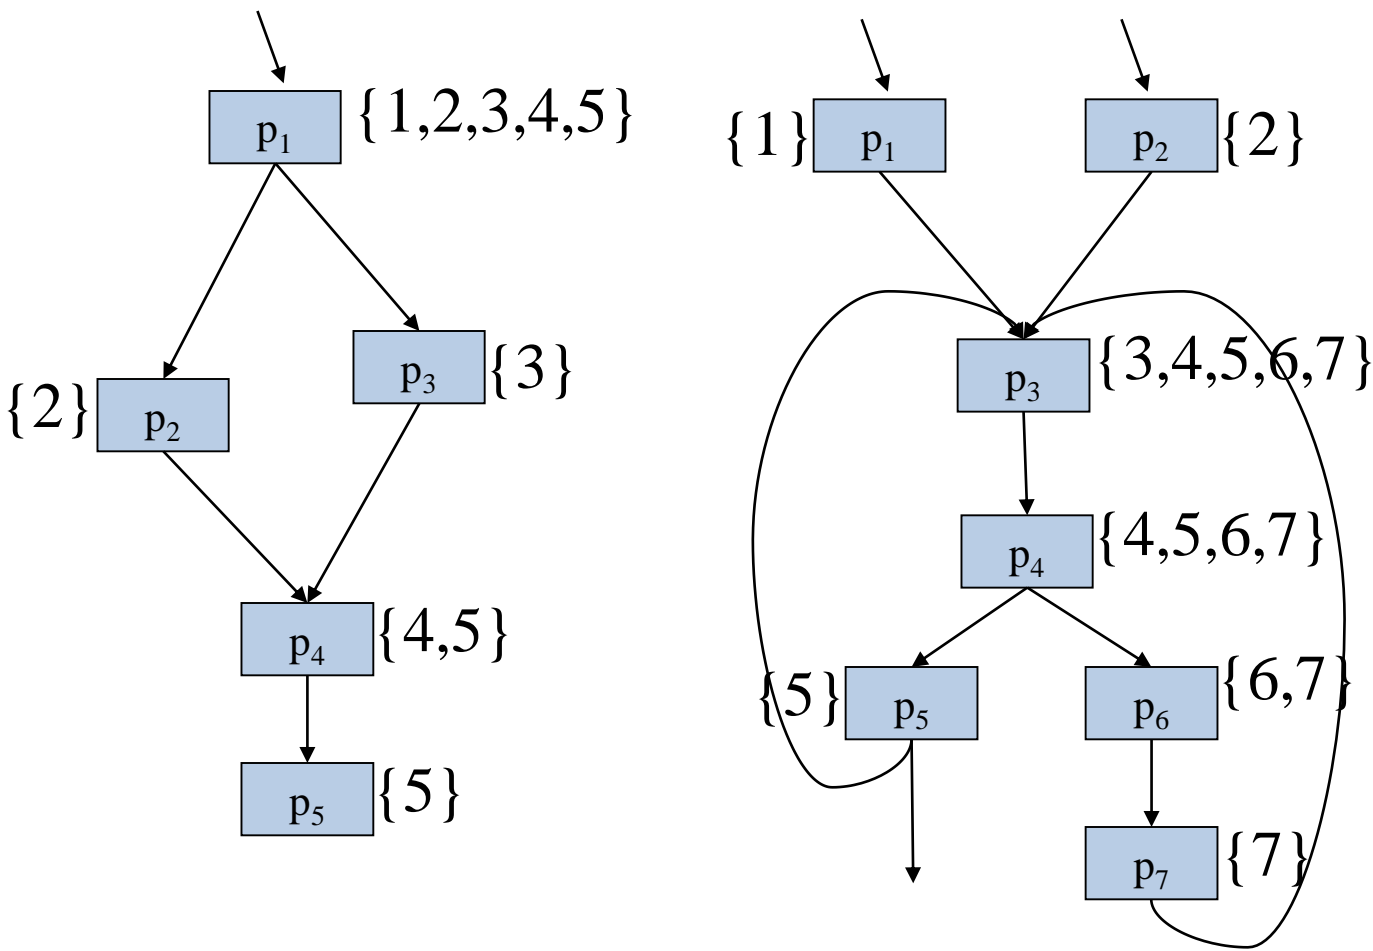
A *loop* in a control flow graph is a set of nodes S including a *header* node h , with the following properties:

 - From any node in S there is a path leading to h
 - There is a path from h to any node in S
 - There is no edge from any node outside S to any node in S other than h



- **Definition: dominator**

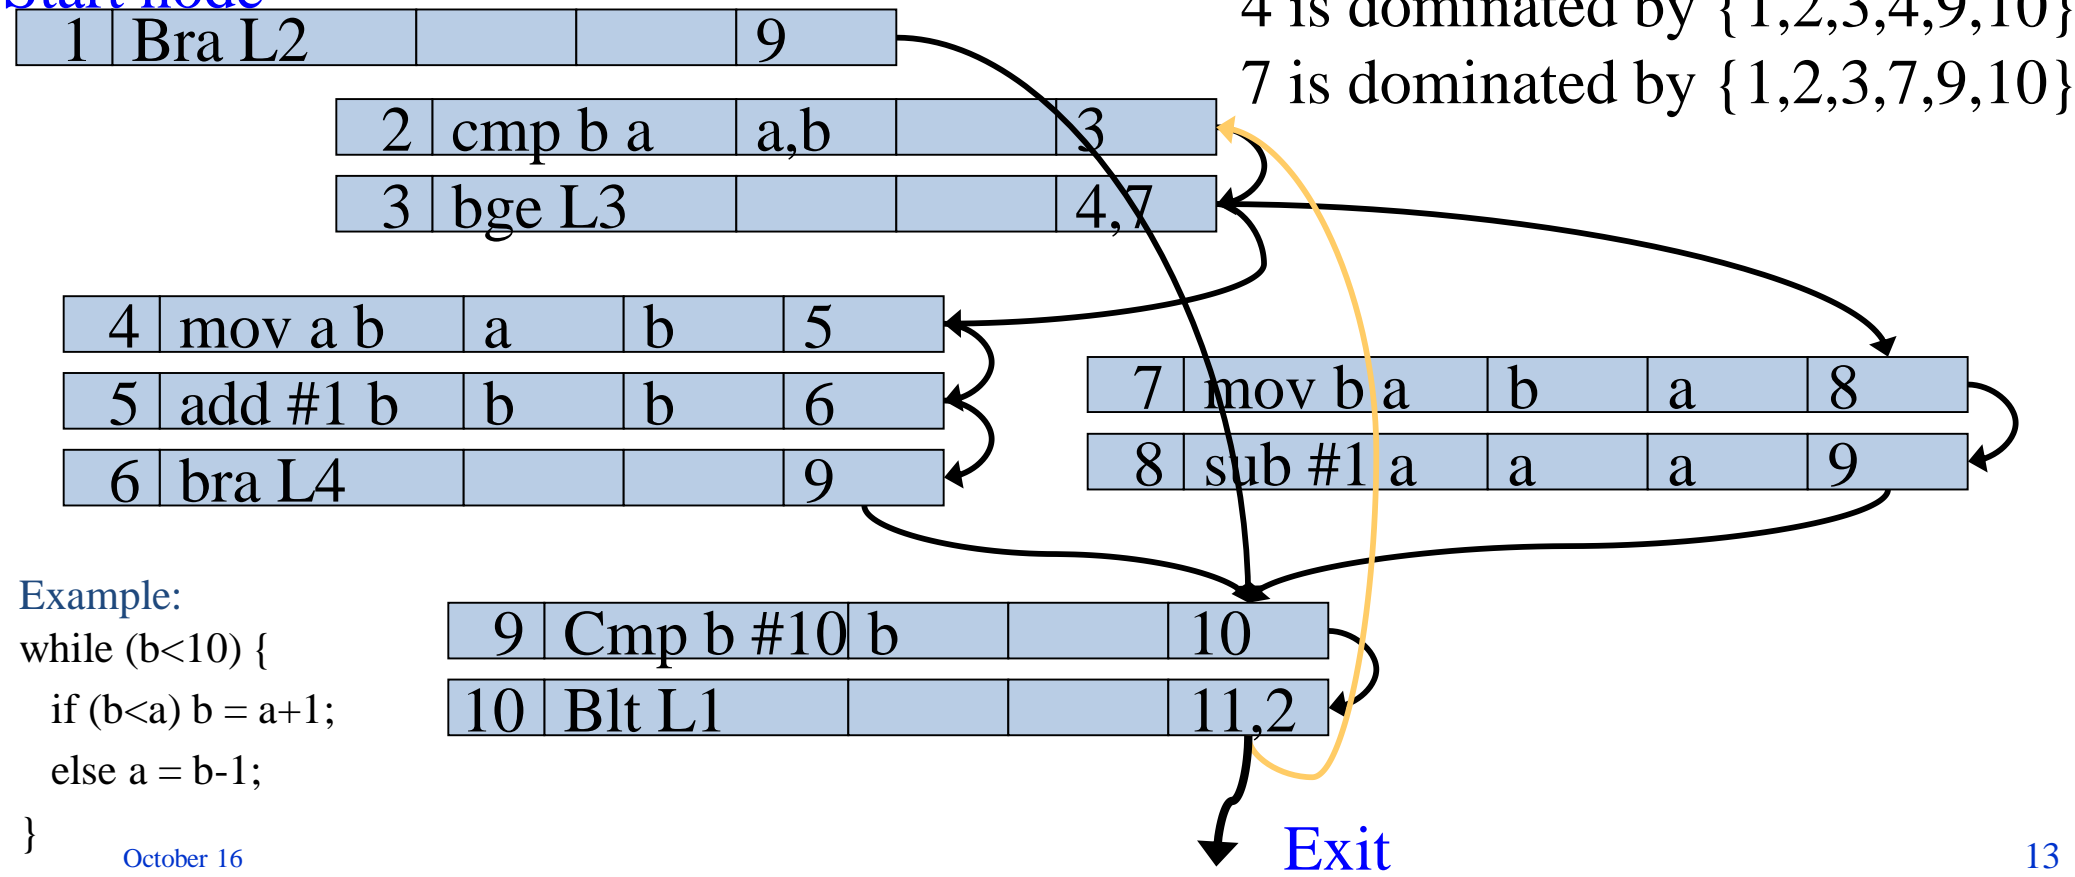
A node d *dominates* a node n if every path from the CFG's start node to n must go through d . Every node dominates itself



• Definition: dominator


A node d *dominates* a node n if every path from the CFG's start node to n must go through d . Every node dominates itself

Start node



Dominators...

- Finding the nodes dominated by a node d :
 - Consider another node n with predecessors $p_1 \dots p_k$
 - If d dominates each one of the p_i then it must dominate n
 - Because:
 - Every path from the start node to n must go through one of the p_i
 - And every path from the start node to a p_i must go through d
 - Conversely,
 - If d dominates n , it must dominate all the p_i
 - Otherwise there would be a path from the start node to n going through the predecessor not dominated by d




Algorithm for finding dominators

- Let $\text{Doms}(n)$ be the set of nodes that dominate n
(*“ n is dominated by $\text{Doms}(n)$ ”*)
- Construct a system of simultaneous set equations:
- $\text{Doms}(s) = \{ s \}$ (*$s = \text{start node}$*)
- $\text{Doms}(n) = \{ n \} \cup \left(\bigcap_{p \in \text{preds}(n)} \text{Doms}(p) \right)$ (*otherwise*)
(*“which dominators are common to all our preds?”*)
- Solve this system iteratively
- Initially, each $\text{Doms}(n)$ starts as the set of all nodes in the graph
- Each assignment makes $\text{Doms}(n)$ smaller, until it stops changing

Back edges


- A control flow graph edge from a node n to a node h that dominates n is called a *back edge*.



n_1 dominates all nodes
 n_2 dominates n_2, n_4
 n_3 dominates only n_3
 n_4 dominates only n_4
 n_5 dominates only n_5
 n_6 dominates only n_6

Back edges...

- For every back edge, there is a corresponding subgraph of the CFG that is a loop (by our definition earlier)



Definition:

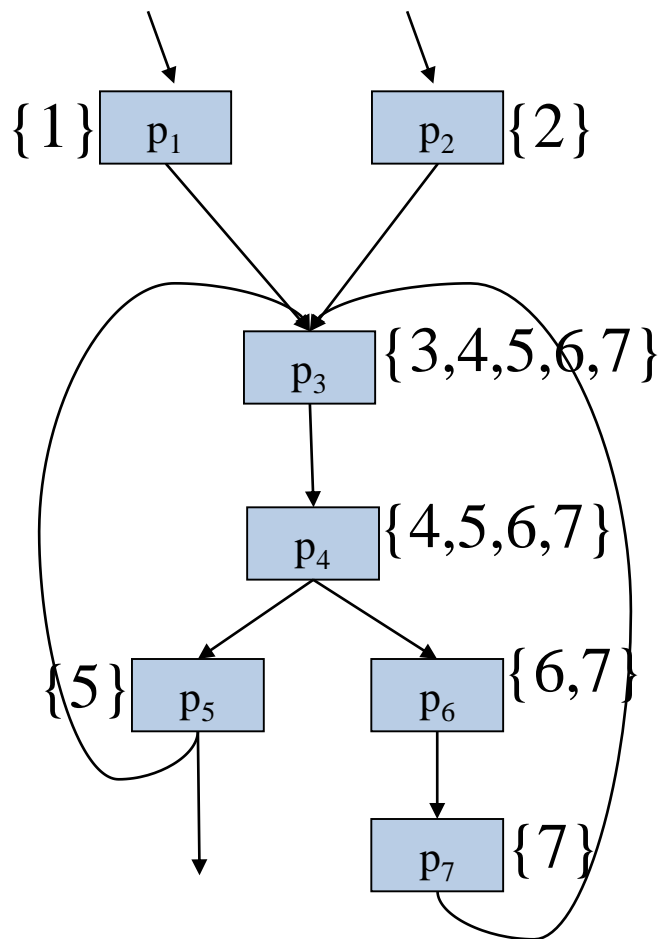
The *natural loop* of a backedge (n, h) , where h dominates n , is

- the set of nodes x such that h dominates x and
- there is a path from x to n not containing h .

The *header* of this loop will be h

Back edges...

- For every back edge, there is a corresponding subgraph of the CFG that is a loop (by our definition earlier)



Definition:


The *natural loop* of a backedge (n,h) , where h dominates n , is

- the set of nodes x such that h dominates x and
- there is a path from x to n not containing h .

The *header* of this loop will be h


Multiple loops

- It is possible for two loops to share the same header
- This example has two back edges, (5,3) and (7,3)
- The easiest thing to do in this case is to treat them as one loop




Multiple loops

- It is possible for two loops to share the same header
- This example has two back edges, (5,3) and (7,3)
- The easiest thing to do in this case is to treat them as one loop



Multiple loops

- It is possible for two loops to share the same header
- This example has two back edges, (5,3) and (7,3)
- The easiest thing to do in this case is to treat them as one loop



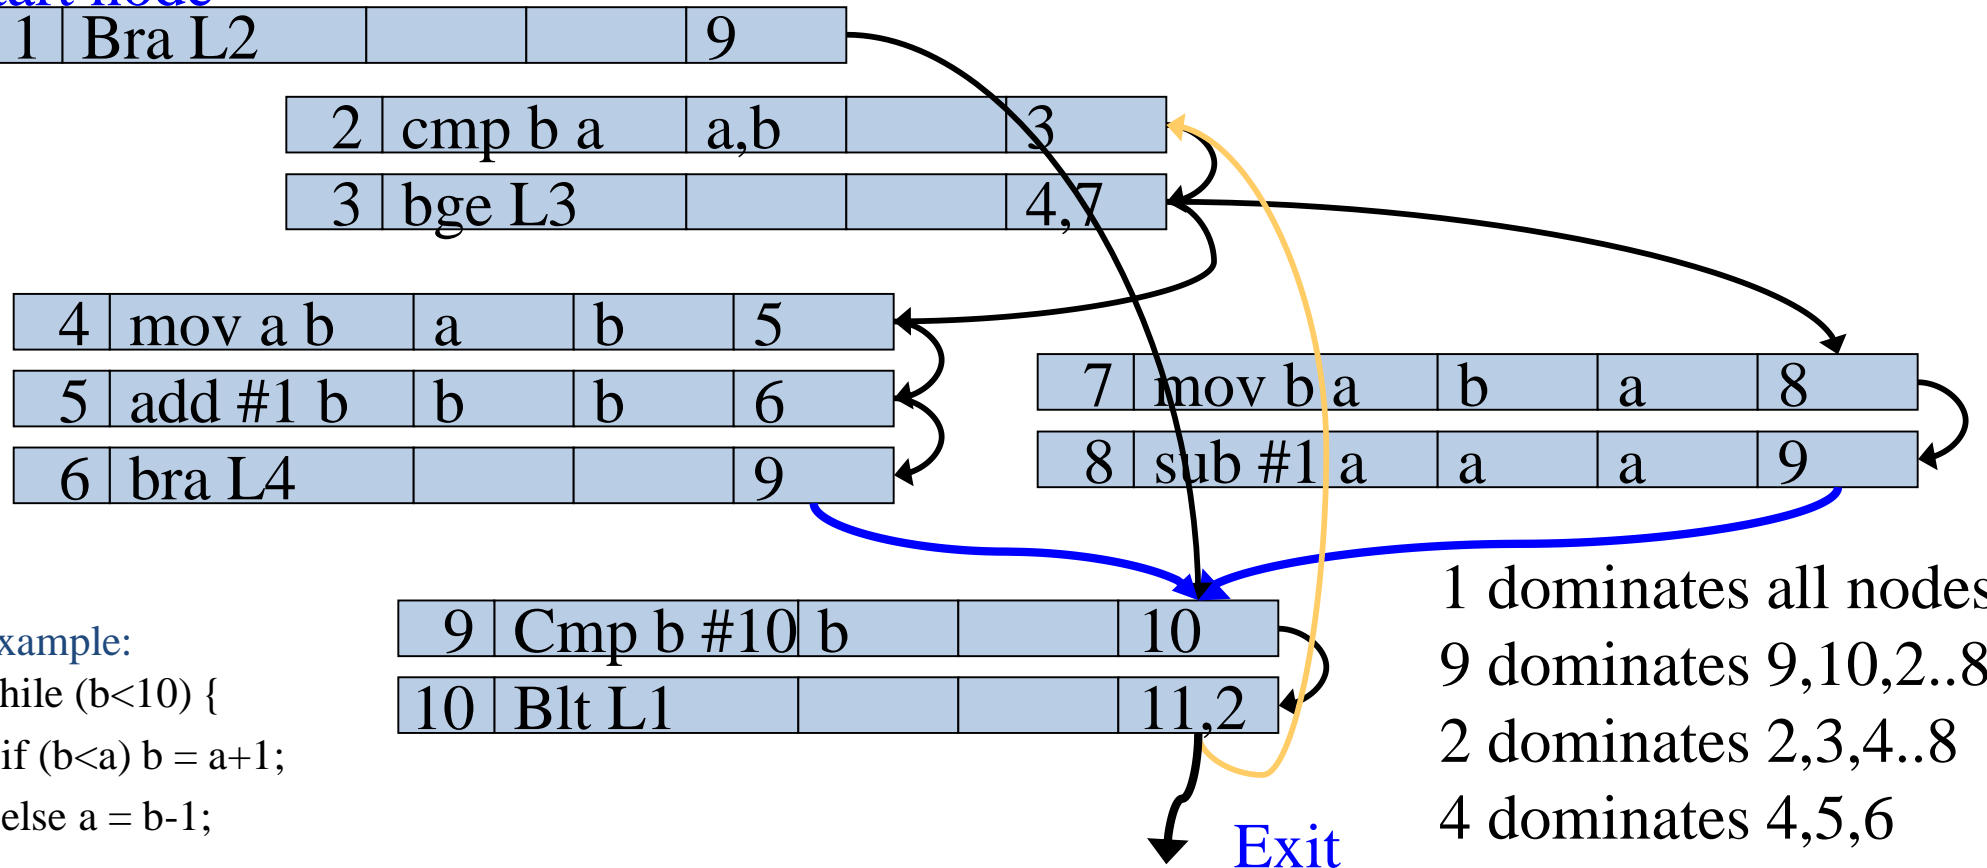
Natural loop of (5,3)

Natural loop of (7,3)

Multiple loops

- It is possible for two loops to share the same header
- This example has two back edges, (6,9) and (8,9)
- The easiest thing to do in this case is to treat them as one loop


Start node



1 dominates all nodes
 9 dominates 9,10,2..8
 2 dominates 2,3,4..8
 4 dominates 4,5,6
 7 dominates 7,8

Nested loops


- Suppose:
 - A and B are loops with headers a and b, such that $a \neq b$, and b is in A
- Then
 - The nodes of B must be a proper subset of the nodes of A
 - We say that loop B is nested within A
 - B is the inner loop



Back edges: (3,2), (4,2), (10,5), (9,8)

Nested loops



- Suppose:
 - A and B are loops with headers a and b, such that $a \neq b$, and b is in A
- Then
 - The nodes of B must be a proper subset of the nodes of A
 - We say that loop B is nested within A
 - B is the inner loop



Back edges: (3,2), (4,2), (10,5), (9,8)

The Control Tree

- Loops form a tree
- Example:





Back edges: (3,2), (4,2), (10,5), (9,8)

We have reconstructed the “structured control flow” from the control flow graph

Pre-headers


- Where should we move the loop-invariant instructions *to*?



- We can't move them to the header
- We want to move them to the node preceding the header


Pre-headers

- Where should we move the loop-invariant instructions *to*?
- We want to move them to the node preceding the header
- But sometimes the header has multiple predecessors
- What shall we do?



Pre-headers

- Where should we move the loop-invariant instructions *to*?
- We want to move them to the node preceding the header
- But sometimes the header has multiple predecessors
- What shall we do?
 - **Insert a pre-header**



Which instructions can we move out of a loop?

- The next question is exactly which loop-invariant instructions we can move to the pre-header
- It's easy to get it wrong....

A

L₀:
t = 0

L₁:
i = i+1
t = a ⊕ b
M[i] = t
if i < N goto L₁

L₂:
x = t

B

L₀:
t = 0

L₁:
if i < N goto L₂
i = i+1
t = a ⊕ b
M[i] = t
goto L₁

L₂:
x = t

Which instructions can we move out of a loop?

- The next question is exactly which loop-invariant instructions we can move to the pre-header
- It's easy to get it wrong....

A

L₀:
t = 0
t = a ⊕ b
L₁:
i = i+1
~~t = a ⊕ b~~
M[i] = t
if i < N goto L₁
L₂:
x = t

B

L₀:
t = 0
t = a ⊕ b
L₁:
if i < N goto L₂
i = i+1
~~t = a ⊕ b~~
M[i] = t
goto L₁
L₂:
x = t

Which instructions can we move out of a loop?

- It's easy to get it wrong....

C

$L_0:$
 $t = 0$

$L_1:$
 $i = i + 1$
 $t = a \oplus b$
 $M[i] = t$
 $t = 0$
 $M[j] = t$
if $i < N$ goto L_1
 $L_2:$

D

$L_0:$
 $t = 0$

$L_1:$
 $M[j] = t$
 $i = i + 1$
 $t = a \oplus b$
 $M[i] = t$
if $i < N$ goto L_1
 $L_2:$
 $x = t$

Which instructions can we move out of a loop?

- It's easy to get it wrong....

C

L_0 :
t = 0
t = a \oplus b

L_1 :
i = i+1
~~t = a \oplus b~~
M[i] = t
t = 0
M[j] = t
if i < N goto L_1

L_2 :

D

L_0 :
t = 0
t = a \oplus b

L_1 :
M[j] = t
i = i+1
~~t = a \oplus b~~
M[i] = t
if i < N goto L_1

L_2 :

Which instructions can we move out of a loop?

<p>A</p> <p>L_0: $t = 0$ $t = a \oplus b$</p> <p>L_1: $i = i + 1$ $t = a \oplus b$ $M[i] = t$ if $i < N$ goto L_1</p> <p>L_2: $x = t$</p>	<p>B</p> <p>L_0: $t = 0$ $t = a \oplus b$</p> <p>L_1: if $i < N$ goto L_2 $i = i + 1$ $t = a \oplus b$ $M[i] = t$ goto L_1</p> <p>L_2: $x = t$</p>	<p>C</p> <p>L_0: $t = 0$ $t = a \oplus b$</p> <p>L_1: $i = i + 1$ $t = a \oplus b$ $M[i] = t$ $t = 0$ $M[j] = t$ if $i < N$ goto L_1</p> <p>L_2:</p>	<p>D</p> <p>L_0: $t = 0$ $t = a \oplus b$</p> <p>L_1: $M[j] = t$ $i = i + 1$ $t = a \oplus b$ $M[i] = t$ if $i < N$ goto L_1</p> <p>L_2:</p>
---	---	---	--

Hoist

Don't hoist:

Loop invariant node
does not dominate
all loop exits

Don't hoist:

More than one
definition of t
in the loop

Don't hoist:

t is liveOut from
the loop's
preheader

Which instructions can we move out of a loop?

- Conditions for hoisting a CFG node

d: $t = a \oplus b$

1 Loop invariant: all reaching defs used by d occur outside loop

Use Reaching Definitions data flow analysis

2 Loop invariant node must dominate all loop exits

Use Dominators analysis

3 There must be just one def of t in loop

Just count them!

4 t must not be liveOut from the loop's preheader

Use Live Variables data flow analysis

What next...

- Hoisting loop invariants really helps
- But good compilers do lots more...
 - Induction variables:
 - A variable which is incremented by a loop-invariant amount
 - A variable which is a multiple of an induction variable
 - Strength reduction
 - Compute all induction variables by incrementing instead of multiplying
 - Induction variable elimination, rewriting comparisons
 - Array bounds check elimination
 - Range of all induction variables is known on entry to a for loop
 - Common sub-expressions
 - More sophisticated methods – eg partial redundancy elimination
- Now you have seen how to hoist loop-invariants, you can figure the rest out yourself!

Optimisations for high-performance computing

- “Conventional” optimisations *reduce* work done at run-time
- “restructuring” compilers improve performance by finding the *right order* in which to do the computation
- Example: Parallelisation:

Original code:

```
For (i=0;i<N;i++)
```

```
  For (j=0;j<M;j++)
```

```
    A[i,j] = (A[i,j] + A[i-1,j] + A[i+1,j])* (1/3)
```

Parallel implementation:

```
For (i=0;i<N;i++)
```

```
  ParFor (j=0;j<M;j++)
```

```
    A[i,j] = (A[i,j] + A[i-1,j] + A[i+1,j])* (1/3)
```

Better parallel implementation?

```
ParFor (j=0;j<M;j++)
```

```
  For (i=0;i<N;i++)
```

```
    A[i,j] = (A[i,j] + A[i-1,j] + A[i+1,j])* (1/3)
```

Optimisations for high-performance computing

- Another restructuring example:

Example: matrix transpose:

```
for (i=0;i<N;i++)  
    for (j=0;j<M;j++)  
        B[i][j] = A[j][i];
```

Cache-efficient implementation:

```
for (ii=0;ii<N;ii+=IB)  
    for (jj=0;jj<M;jj+=JB)  
        for (i=ii;i<ii+IB;i++)  
            for (j=jj;j<jj+JB;j++)  
                B[i][j] = A[j][i];
```

Using 1.4GHz AMD Athlon
N=M=2000
IB=JB=100
Original execution time: 462ms
Improved execution time: 96ms

Optimisations for high-level programming languages

- Subtype polymorphism
 - Static resolution of the type of x in $x.f()$ enables inlining of method f
- Generics (aka parametric polymorphism)
 - A generic class is parameterised by a type (eg a container by its element type). When is it a good idea to generate specialised code?
- Pattern matching
 - In a language like Haskell or Prolog, pattern matching on nested data structures is very powerful. Find optimum sequence of tests.
- Dynamic object creation
 - If we allow space to be allocated, but automatically freed, can we sometimes add code to do it instead of relying on garbage collection?
- Lazy evaluation
 - Can an expression be evaluated where it is first referred to, or do we have to build a “closure” representing it?
- Arrays – overloaded arithmetic
 - If we overload arithmetic operators to work on arrays, how to avoid lots of little loops?
- Arrays – slices
 - If we allow a multidimensional array to be sliced, eg $A[2:99,2:99]$, how do we avoid having to manipulate an array descriptor?

Textbooks

EaC

- Data flow analysis is covered in Chapter 9
 - Reaching definitions are covered in Section 9.2.4
 - Dominators are covered in Section 9.3.2
- EaC handles loop-invariant code motion somewhat differently from these slides, which are based on Appel's presentation
 - See “Lazy Code Motion” (LCM), page 506
 - LCM resolves the hoisting conditions in a more systematic way than presented here, by combining four different data-flow analyses

Textbooks

- Appel also covers optimisation in depth
 - Chapter 10 introduces DFA through live variable analysis
 - Chapter 17 shows how DFA can be used for many other useful analyses
 - Chapter 18 deals with finding loops, finding induction variables, and implementing loop optimisations (which rely on DFAs)
 - Chapter 19 presents Static Single Assignment, a program representation which provides easy (and space-efficient) access to dependence information such as reaching definitions. This simplifies many loop optimisations
 - Chapter 20 covers instruction scheduling – finding an instruction ordering which makes optimal use of modern CPU architectures
 - Chapter 21 concerns improving cache performance – by prefetching, and by executing loops blockwise
- Another really good source if you're building an optimising compiler is “High-performance compilers for parallel computing”, Michael Wolfe (Addison Wesley 1996)
- Fine print:
 - CFG would consist of basic blocks instead of individual instructions
 - For loop optimisations, we would do the DFA on the IR before instruction selection; it's simpler and it avoids complications such as having only two-address instructions
 - See Appel pg388
- Credits: in addition to Appel's book, I found it very useful to study the course notes of Liz White (George Mason University), Laurie Hendren (McGill University) and Chau-Wen Tseng (University of Maryland)

Research

- Several Imperial research groups are working on optimising compiler technology, including:
 - Wayne Luk's Custom Computing/Silicon Compilation group
 - Alastair Donaldson's group
 - Paul Kelly's Software Performance Optimisation group
 - Compiler-related research: Cristian Cadar, Peter Pietzuch, Sergio Maffeis etc
 - Programming languages: Sophia Drossopoulou, Nobuko Yoshida, and others
- Opportunities: UROP summer placements, individual projects, and PhDs
- Sample projects:
 - Automatically searching for the best combination of blocking, loop fusion, unrolling, parallelisation, vectorisation
 - Work computational scientists to make their simulation of tidal turbines/Formula 1/blood flow/weather run fast on 10,000-100,000 cores
 - Efficient execution of analysis queries on results from large parallel fluid dynamics simulations
 - Design a domain-specific language and compiler to generate high-performance code for 3D robot vision and scene understanding

Implementing loop optimisations in Haskell

- The next few slides give a Haskell implementation for some of the ideas presented in this chapter
- This material is provided to provide a concrete illustration of the concepts
- It is the concepts which are important, not the code
- **Do *not* memorise the code** – spend the time reading the textbook instead
- Some of the algorithms used here are rather inefficient – in many cases we just transcribe the mathematical definitions. Efficient algorithms exist – but are considerably more complicated.

Reaching definitions – gen and kill

- Preliminaries: the Gen and Kill sets:

```
nodeGen node | nodeDefs node == [] = []  
           | otherwise      = [nodeId node]  
  
nodeKill cfg node = nodeDefSet cfg node \\ [nodeId node]
```

- Suppose t is defined in $node$. $nodeDefSet$ is set of all the nodeids where t is defined:

```
nodeDefSet (ControlFlowGraph cfg) node  
= case nodeDefs node of  
  [t] -> [id | Node id i ds us scs prds <- cfg,  
            t `elem` ds]  
  []   -> []  
  otherwise -> error "nodeDefSet: multiple defs"
```

- Auxiliary functions used in solver overleaf:

```
untilConverges (a:b:rest) | a == b = a  
untilConverges (a:b:rest)      = untilConverges (b:rest)  
  
zip2 (rdsin,rdsout) = zip rdsin rdsout  
bigU sets = nub (concat sets)
```

Reaching definitions - solver

- Solve the dataflow equations:

```
reachingDefinitionsOf :: CFG -> ( [ (Id,[Id]) ], [ (Id,[Id]) ] )

reachingDefinitionsOf cfg
= untilConverges (iterate updateRDs initialRDs)
  where
    initialRDs :: ( [ (Id,[Id]) ], [ (Id,[Id]) ] )
    initialRDs = ( [(n,[])] | n<-nodesOf cfg, [(n,[])] | n<-nodesOf cfg )

    updateRDs :: ( [(Id,[Id])], [(Id,[Id])] ) -> ( [(Id,[Id])], [(Id,[Id])] )
    updateRDs rds = unzip (map (updateRD rds) (zip2 rds))
    updateRD (rdins_sofar,rdouts_sofar) ((id,rdins), (sameid,rdouts))
      = ((id,rdins'), (id,rdouts'))
      where
        rdins' = bigU [retrieve s rdouts_sofar | s <- nodePreds node]
        rdouts' = nodeGen node `union` ((rdInsOf node) \\ nodeKill cfg node)
        where
          rdInsOf node = retrieve (nodeId node) rdins_sofar
          node = idToNode cfg id
```

- We solve the system of simultaneous set equations iteratively
- Initially each node's ReachIn (rdins), and ReachOut (rdouts) set is empty
- The updates successively increase the ReachIn and ReachOut sets until convergence

Use reaching definitions to find loop invariant instructions

- Find the definitions which reach this node which are relevant – that is, which generate the values this node uses:

```
relevantReachingDefinitionsOf :: CFG -> [ (Id,[Id]) ]
```

```
relevantReachingDefinitionsOf cfg
```

```
= [(nodeId node, relevantDefs node) | node <- cfgToNodes cfg]
```

```
  where
```

```
    relevantDefs node
```

```
      = [rd | rd <- retrieve (nodeId node) rds_in,
```

```
            nodeDefs (idToNode cfg rd) `intersect` nodeUses node /= []]
```

```
    (rds_in, rds_out) = reachingDefinitionsOf cfg
```

Use reaching definitions to find loop invariant instructions

- An instruction is loop invariant if the definitions of all the values it uses are outside the loop:

```
> externallyDependentInstructionsOf cfg loop
> = [node | node <- [idToNode cfg id | id <- loop],
>     nodeDefs node /= [],
>     relevantDefs node `intersect` loop == [],
>     hoistable node]
> where
> relevantDefs node = retrieve (nodeId node)
>                       (relevantReachingDefinitionsOf cfg)
```

- An instruction is hoistable only if it produces a value (ie not a compare, branch, etc):

```
hoistable (Node id i [] uses succs preds) = False
hoistable (Node id i defs uses succs preds) = True
```

Use reaching definitions to find loop invariant instructions

- Now iteratively add instructions which are l-i because they depend only on l-i instructions. We reverse the result so that when we add them to the pre-header, they are added in dependence-order.

```
> loopInvariantInstructionsOf cfg loop
> = reverse (untilConverges (iterate updateLIs initialLIs))
>   where
>     initialLIs = externallyDependentInstructionsOf cfg loop
>     updateLIs :: [CFGNode] -> [CFGNode]
>     updateLIs invariantsSoFar
>       = invariantsSoFar `union`
>         [n | n <- map (idToNode cfg) loop,
>           hoistable n,
>           and [hasSingleInvariantDefinition n u | u<-nodeUses n]]
>     where
>       hasSingleInvariantDefinition n u
>         = length defs == 1 && head defs `elem` map nodeId invariantsSoFar
>         where
>           defs = [d | d<-relevantDefs n, u `elem` nodeDefs (idToNode cfg d)]
```

Finding dominators... implementation

```
dominatorsOf :: CFG -> [(Id,[Id])]
dominatorsOf cfg
= untilConverges (iterate updateDs initialDs)
  where
    initialDs :: [(Id,[Id])]
    initialDs = [ (n, nodesOf cfg) | n <- (nodesOf cfg)]

    updateDs :: [(Id,[Id])] -> [(Id,[Id])]
    updateD ds_sofar (id,d)
      = (id,
          [id] `union` (bigCap [retrieve p ds_sofar | p <- nodePredsOf id])
        )
    updateDs ds = map (updateD ds) ds

    nodePredsOf id = nodePreds (idToNode cfg id)

    bigCap [] = []
    bigCap sets = foldr1 intersect sets

    untilConverges (a:b:rest) | a == b = a
    untilConverges (a:b:rest)          = untilConverges (b:rest)
```

- We solve the system of simultaneous set equations iteratively
- Initially each node's Doms set is the set of all the nodes of the CFG
- The updates successively reduce the Doms until convergence

Finding back edges

- A flow graph edge from a node n to a node h that dominates n is called a back edge:

```
backEdges :: CFG -> [(Id,Id)]
```

```
backEdges cfg
```

```
= [ (n,h) | n <- nodesOf cfg, h <- nodesOf cfg, n /= h,  
        flowedge n h,  
        h `dominates` n]
```

where

```
dominators = dominatorsOf cfg
```

```
a `dominates` b = a `elem` (retrieve b dominators)
```

```
flowedge a b = a `elem` nodePreds (idToNode cfg b)
```

Finding natural loops

- The *natural loop* of a backedge (n,h), where h dominates n, is the set of nodes x such that h dominates x and there is a path from x to n not containing h.

```
naturalLoop :: CFG -> (Id,Id) -> (Id, [Id])
--                backedge  header, nodes

naturalLoop cfg (n,header)
= (header, real_xs)
  where
    poss_xs = [x | x <- nodesOf cfg, header `dominates` x]
    real_xs = [x | x <- poss_xs, pathExists x n]
    pathExists x n
      = [] /= [path | path <- allpaths, not (header `elem` path)]
      where
        allpaths = findControlFlowPaths cfg x n
        dominators = dominatorsOf cfg
        a `dominates` b = a `elem` (retrieve b dominators)
```

(omit paths via header, and therefore paths via enclosing loops)

(findControlFlowPaths defined next slide)

Finding paths

- I have used a general-purpose path enumeration to find all the paths from one node to another. This is rather wasteful... Some care is needed to avoid following cycles; "mypath" below records the nodes visited so far.

```
findControlFlowPaths :: CFG -> Id -> Id -> [[Id]]
```

```
findControlFlowPaths cfg start end = findControlFlowPaths' [] start
```

```
  where
```

```
  findControlFlowPaths' mypath x
```

```
    | x == end          = [[x]]
```

```
    | x `elem` mypath    = [[]]
```

```
    | otherwise         = map (x:) restOfPath
```

```
      where
```

```
      extendedpath = x:mypath
```

```
      succs = nodeSuccs (idToNode cfg x)
```

```
      nonCycleSuccs = succs
```

```
      restOfPath = concat (map (findControlFlowPaths' extendedpath) nonCycleSuccs)
```

Building the loop nest tree (a.k.a. the control tree)

- The loop nest tree consists at each level of a loop (with its header), and the list of all its subloop trees:

```
data LoopTree = LTree (Id,[Id]) [LoopTree] deriving (Show, Eq)
```

```
loopTree :: CFG -> LoopTree
```

```
loopTree cfg
```

```
  = LTree (0, nodesOf cfg) (makeTrees theloops)
```

```
  where
```

```
    backedges = backEdges cfg
```

```
    theloops = map (naturalLoop cfg) backedges
```

```
    makeTrees loops = map makeTree (siblingloops loops)
```

```
    makeTree loop
```

```
      = LTree loop (makeTrees subloops)
```

```
      where
```

```
        subloops = [(h,nub l) | (h,l) <- theloops, containedIn (h,l) loop]
```

Building the loop nest tree...

- The children of a given loop are the immediate subloops. A subloop is an immediate subloop if it is not contained in any other loop in the list:

siblingloops loops

= [l1 | l1 <- loops,

not (any (containedIn l1) [l2 | l2 <- loops, l1 /= l2])]

- To work out whether one loop l1 is strictly contained within another l2, we ask simply whether l1's header is in l2's body:

containedIn :: (Id,[Id]) -> (Id,[Id]) -> Bool

containedIn (h1,l1) (h2,l2) = h1 `elem` l2

Manipulating the control flow graph...

- To implement hoisting of loop invariants we need a few other functions:
 - Insert a pre-header before each loop header:
 - > addPreHeaders :: CFG -> LoopTree -> [(Id,Id)], CFG)
 - > addPreHeaders cfg looptree =
 - Remove a specified list of nodes from a cfg
 - > removeNodes :: [CFGNode] -> CFG -> CFG
 - > removeNode node cfg = ...
 - Insert a specified node n into a cfg after a specified node "target". This only works if the target has only one successor, as is the case with a pre-header.
 - > [CFGNode] -> CFG -> Int -> CFG
 - > insertNodesAfter nodes cfg target = ...
 - Traverse the modified CFG and generate instructions:
 - > generateInstructions :: CFG -> [Instruction]
 - > generateInstructions cfg = ...

Hoisting the loop-invariants...

- Finally, we bring it all together
 - > hoistLoopInvariants cfg looptree
 - > = newcfg
 - > where
 - > newcfg = foldl hoistALoop cfgWithPreheaders loops
 - > loops = [(h,l) | (h,l) <- loopsOf looptree, h /= 0]
 - > (preheaders, cfgWithPreheaders) = addPreHeaders cfg looptree
 - > hoistALoop cfg (header,body)
 - > = insertNodesAfter invariants (removeNodes invariants cfg) preheader
 - > where
 - > invariants = loopInvariantInstructionsOf cfg (header:body)
 - > preheader = retrieve header preheaders
 - > loopsOf (LTree (h,body) subloops)
 - > = (h,body) : concat (map loopsOf subloops)

*(This sketch
implementation
doesn't check all the
hoisting conditions...)*

AST	Original control flow graph:
(Program	Node 0 (Mov (ImmNum 1) (Reg T1)) [T1] [] [1] []
[(Decl "w" Integer),	Node 1 (Mov (ImmNum 100) (Reg T0)) [T0] [] [2] [0]
(Decl "x" Integer),	Node 2 (Mov (ImmNum 200) (Reg T3)) [T3] [] [3] [1]
(Decl "y" Integer),	Node 3 (Mov (Reg T1) (Reg T4)) [T4] [T1] [4] [2,15]
(Decl "z" Integer)]	Node 4 (Add (ImmNum 1) (Reg T4)) [T4] [T4] [5] [3]
[Assign (Var "x") (Const 1),	Node 5 (Mov (Reg T4) (Reg T1)) [T1] [T4] [6] [4]
Assign (Var "w") (Const 100),	Node 6 (Mov (Reg T3) (Reg T5)) [T5] [T3] [7] [5]
Assign (Var "z") (Const 200),	Node 7 (Mov (Reg T0) (Reg T6)) [T6] [T0] [8] [6]
LabelStat "Here",	Node 8 (Add (Reg T5) (Reg T6)) [T6] [T5,T6] [9] [7]
Assign (Var "x") (Binop Plus	Node 9 (Mov (Reg T6) (Reg T2)) [T2] [T6] [10] [8]
(Ref (Var "x")) (Const 1)),	Node 10 (Mov (Reg T1) (Reg T7)) [T7] [T1] [11] [9]
Assign (Var "y") (Binop Plus	Node 11 (Mov (ImmNum 10) (Reg T8)) [T8] [] [12] [10]
(Ref (Var "w")) (Ref (Var "z"))),	Node 12 (Cmp (Reg T7) (Reg T8)) [] [T7,T8] [13] [11]
IfThenElse (Compare CLT	Node 13 (Blt "L1") [] [] [14,15] [12]
(Ref (Var "x")) (Const 10))	Node 14 (Bra "L2") [] [] [17] [13]
[Goto "Here"] []	Node 15 (Bra "LHere") [] [] [3] [13]
])	Node 16 (Bra "L3") [] [] [17] []
	Node 17 Halt [] [] [] [14,16]

- Relevant reaching definitions:
relevantReachingDefinitionsOf cfg =

```

(0,[]),
(1,[]),
(2,[]),
(3,[0,5]),
(4,[3]),
(5,[4]),
(6,[2]),
(7,[1]),
(8,[7,6]),
(9,[8]),
(10,[5]),
(11,[]),
(12,[11,10]),
(13,[]),
(14,[]),
(15,[]),
(16,[]), (17,[])]

```

- Loop Tree:

```

loopTree cfg =
    LTree (0,[0,1,2,3,4,5,6,7,8,9,10,
              11,12,13,14,15,16,17])
    [LTree (3,[4,5,6,7,8,9,10,
              11,12,13,15])
      [] ]

```

- Loop invariants:

```

externallyDependentInstructionsOf cfg
  [3,4,5,6,7,8,9,10,11,12,13,15] =
  [Node 6 (Mov (Reg T3) (Reg T5)) [T5] [T3] [7] [5],
   Node 7 (Mov (Reg T0) (Reg T6)) [T6] [T0] [8] [6],
   Node 11 (Mov (ImmNum 10) (Reg T8)) [T8] [] [12] [10]]

```

```

loopInvariantInstructionsOf (cfgex 15)
  [3,4,5,6,7,8,9,10,11,12,13,15] =
  [Node 7 (Mov (Reg T0) (Reg T6)) [T6] [T0] [8] [6],
   Node 6 (Mov (Reg T3) (Reg T5)) [T5] [T3] [7] [5],
   Node 11 (Mov (ImmNum 10) (Reg T8)) [T8] [] [12] [10],
   Node 9 (Mov (Reg T6) (Reg T2)) [T2] [T6] [10] [8],
   Node 8 (Add (Reg T5) (Reg T6)) [T6] [T5,T6] [9] [7]]

```

- Code after loop-invariant hoisting:

```
move.l #1, T1
move.l #100, T0
move.l #200, T3
#Preheader for loop with header 3
move.l T0, T6
move.l T3, T5
move.l #10, T8
move.l T6, T2
add.l T5, T6
```

(continued in next column...)

```
M3:
move.l T1, T4
add.l #1, T4
move.l T4, T1
#Mov (Reg T3) (Reg T5) moved
#Mov (Reg T0) (Reg T6) moved
#Add (Reg T5) (Reg T6) moved
#Mov (Reg T6) (Reg T2) moved
move.l T1, T7
#Mov (ImmNum 10) (Reg T8) moved
cmp.l T7, T8
blt M15
bra M14

M15:
bra M3

M14:
bra M17

M17:
halt
bra M3
```

- Fine print:
 - For efficiency, it is better for the CFG to consist of basic blocks instead of individual instructions
 - For loop optimisations, we would do the DFA on the IR before instruction selection; it's simpler and it avoids complications such as having only two-address instructions
 - See Appel pg388
- Credits: the primary source for these slides was Appel's book. I also found it very useful to study the course notes of Liz White (George Mason University), Laurie Hendren (McGill University) and Chau-Wen Tseng (University of Maryland)