

Assessed coursework – code generation for functions

Notes and feedback on solutions

Most people did pretty well

There were a few subtle errors...

- Not removing the parameter register when evaluating the actual parameter expression
- Forgetting to reverse pop instructions when restoring registers
- What is the “weight” of “apply”?

Keeping the parameter register safe

```
transExp (Apply s e) (dst:rest)
= (saveRegs ...) ++
  transExp e (paramReg:dst:rest) ++
  [Jsr s] ++
  [Mov (Reg resultReg) (Reg dst)] ++
  (restoreRegs ...)
```

Simple, elegant, wrong
In two ways...

Keeping the parameter register safe

```
transExp (Apply s e) (dst:rest)
= (saveRegs regsInUse) ++
  transExp e (allRegs \\ [paramReg]) ++
  [Mov (Reg (head (allRegs \\ paramReg)) (Reg paramReg))] ++
  [Jsr s] ++
  [Mov (Reg resultReg) (Reg dst)] ++
  (restoreRegs regsInUse)
where regsInUse = allRegs \\ (dst:rest)
```

After
saveRegs,
all registers
are free

Except one

*Once registers
have been
saved, they're
all free*

Push push push, pop pop pop

saveRegs regsInUse

= [Mov (Reg r) Push | r <- regsInUse]

restoreRegs regsInUse

= [Mov Pop (Reg r) | r <- **reverse** regsInUse]

weight of “apply”

- The purpose of the “weight” is to decide the order of subexpression evaluation

- Consider:

$$e1 + f(x)$$

Which is better – $e1$ first? Or do the call first?

weight of “apply”

- The purpose of the “weight” is to decide the order of subexpression evaluation

- Consider:

$e1 + f(x)$

Which is better – $e1$ first? Or do the call first?

If you evaluate $e1$ first, you have to store it in a register

When you call “ f ”, you have to save and restore that register

weight of “apply”

- The purpose of the “weight” is to decide the order of subexpression evaluation

- Consider:

$e1 + f(x)$

Which is better – $e1$ first? Or do the call first?

If you evaluate $e1$ first, you have to store it in a register

When you call “ f ”, you have to save and restore that register

So better to give function calls a high weight

Interprocedural register allocation

- We should be able to avoid unnecessary saving of registers, and we should be able to pass a callee's parameters/results without displacing the caller's.
- Class/file-private functions could be register-allocated at compile-time, but to do it globally it would have to be at link-time.
- It's not common because linking has to be fast but operates on whole programs
- It's also infeasible when linking dynamically.
- It's attractive where linking is static, and inlining is undesirable - eg in space-constrained embedded systems
- And when you have lots of registers

Examples

- SN systems compilers for Sony Playstation 2 and PSP:
 - <http://www.snsys.com/psp/prodg.asp>
- Altium's VX compiler toolset for the ARM processor:
 - <http://www.tasking.com/products/ARM/ARM-ds.pdf>
- Lots of research projects

- Profile-directed:
 - Allocate registers to values used on hot paths
- “leaf” functions are functions that have no callees
 - “About 1/3” of functions are “syntactic” leaf functions
 - “About 2/3” of actual calls are to “effective” leaf functions - no callee is actually called
- Allocate so spill code is on cold paths