

# Operating Systems

Threads

## Processes recap

- One of the **oldest abstractions** in computing
  - An instance of a program being executed, a running program
- Allows a single processor to run multiple programs “simultaneously”
  - Processes turn a single CPU into multiple virtual CPUs
  - Each process runs on a virtual CPU

## Why Have Processes?


- Provide (the illusion of) concurrency
  - Real vs. apparent concurrency
- Provide isolation
  - Each process has its own address space
- Simplicity of programming
  - Firefox doesn't need to worry about gcc
- Allow better utilisation of machine resources
  - Different processes require different resources at a certain time

## Remember optimizing `fork()`

- `fork()`: Creates a new child process by making an exact copy of the parent process image
- Copying the entire address space is expensive!
  - And very few memory pages are going to end up having different values in the two processes



## Page tables: high level view



Schematic view: refer to “Memory management” lectures for details

## Copy On Write (COW)

- Give child its own page table pointing to parent's pages which are marked as *read only*
- When any process writes to a page
  - Protection fault causes trap to the kernel
  - Kernel allocates new copy of page so that both processes have their own private copies
  - Both copies are marked *read-write*



## Dirty Cow Exploit

- Difference between a clean abstraction and a buggy implementation
- <http://thehackernews.com/2016/10/linux-kernel-exploit.html>

# Operating Systems

Threads



## What Are Threads?

- Execution streams that share *the same address space*
- When multithreading is used, each process can contain one or more threads

Per process items	Per thread items
Address space	Program counter (PC)
Global variables	Registers
Open files	Stack
Child processes	
Signals	

## Why Threads?


- Many applications contain multiple activities
  - Which execute in parallel
  - Which access and process the same data
  - Some of which might block

## Example


### Processing thread

- processes input buffer
- writes result into output buffer


Four score and seven years ago, our fathers brought forth upon this continent a new nation: conceived in liberty, and dedicated to the proposition that all men are created equal. Now we are engaged in a great civil war testing whether that	nation, or any nation so conceived and so dedicated, can long endure. We are met on a great battlefield of that war. We have come to dedicate a portion of that field as a final resting place for those who here gave their	lives that this nation might live. It is altogether fitting and proper that we should do this. But, in a larger sense, we cannot dedicate, we cannot consecrate we cannot hallow this ground. The brave men, living and dead,	who struggled here have consecrated it, far above our poor power to add or detract. The world will little note, nor long remember, what we say here, but it can never forget what they did here. It is for us the living, rather, to be dedicated	here to the unfinished work which they who fought here have thus far so nobly advanced. It is rather for us to be here dedicated to the great task remaining before us, that from these honored dead we take increased devotion to that cause for which	they gave the last full measure of devotion, that we have highly resolved that these dead shall not have died in vain that this nation, under God, shall have a new birth of freedom and that government of the people, for the people
---	--	---	---	---	--



Keyboard



Kernel



Disk

### Input thread

- reads data into buffer

### Output thread

- writes output buffer to disk

## Why Not Processes?

- Many applications contain multiple activities
  - Which execute in parallel
  - Which access and process the same data
  - Some of which might block
- Processes are too heavyweight
  - Difficult to communicate between different address spaces
  - An activity that blocks might switch out the entire application
  - Expensive to context switch between activities
  - Expensive to create/destroy activities

## Threads – Problems/Concerns

- Shared address space
  - Memory corruption
    - One thread can write another thread's stack
  - Concurrency bugs
    - Concurrent access to shared data (e.g., global variables)
- Forking
  - What happens on a **fork()** ?
    - Create a new process with the same number of threads
    - Create a new process with a single thread?
- Signals
  - When a signal arrives, which thread should handle it?

# Case Study: PThreads

## PThreads (Posix Threads)

- Defined by IEEE standard 1003.1c
  - Implemented by most UNIX systems

```
#include <pthread.h>
#include <sys/types.h>
```

```
pthread_t      → type representing a thread
pthread_attr_t → type representing the attributes of a thread
```

# Creating Threads

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start_routine)(void*), void *arg);
```

- Creates a new thread
  - The newly created thread is stored in **\*thread**
  - The function returns 0 if thread was successfully created, or error code
- Arguments:
  - **attr** -> specifies thread attributes, can be **NULL** for default attributes
    - Attributes include: minimum stack size, guard size, detached/ joinable, etc.
  - **start\_routine** -> the C function the thread will start execute once created
  - **arg** -> The argument to be passed to start\_routine (of pointer type **void\***). Can be **NULL** if no arguments are to be passed.



# Terminating Threads

```
void pthread_exit(void *value_ptr);
```

- Terminates the thread and makes **value\_ptr** available to any successful join with the terminating thread
- Called implicitly when the thread's start routine returns
  - But not for the initial thread which started `main()`
  - If **main()** terminates before other threads w/o calling **pthread\_exit()**, the entire process is terminated
  - If **pthread\_exit()** is called in **main()** the process continues executing until the last thread terminates (or **exit()** is called)

## PThread Example (1)

```
#include <pthread.h>
#include <stdio.h>

void *thread_work(void *threadid) {
    long id = (long) threadid;
    printf("Thread %ld\n", id);
}

int main (int argc, char *argv[]) {
    pthread_t threads[5];
    long t;
    for (t=0; t<5; t++)
        pthread_create(&threads[t], NULL,
                      thread_work, (void *)t);
}
```

```
$ gcc pt.c -lpthread
$ ./a.out
Thread 0
Thread 1
Thread 2
Thread 3
Thread 4
```

## Passing Arguments to Threads

- What if we want to pass more than one argument to the start routine?
  - Create a structure containing the arguments and pass a pointer to that structure to `pthread_create()`

## Yielding the CPU

```
int pthread_yield(void)
```

- Releases the CPU to let another thread run
- Returns 0 on success, or an error code
  - Always succeeds on Linux
- Why would a thread ever yield? (compare with nice() for processes)

## Joining Other Threads

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- Blocks until **thread** terminates
- The value passed to **pthread\_exit()** by the terminating thread is available in the location referenced by **value\_ptr**
  - **value\_ptr** can be **NULL**

## Join Example

```
#include <pthread.h>
#include <stdio.h>

long a, b, c;
void *work1(void *x) { a = (long)x * (long)x; }
void *work2(void *y) { b = (long)y * (long)y; }

int main (int argc, char *argv[]) {
    pthread_t t1, t2;
    pthread_create(&t1, NULL, work1, (void*) 3);
    pthread_create(&t2, NULL, work2, (void*) 4);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    c = a + b;
    printf("3^2 + 4^2 = %ld\n", c);
}
```

```
$ ./a.out
3^2 + 4^2 = 25
```

## Two Ways to Implement Threads

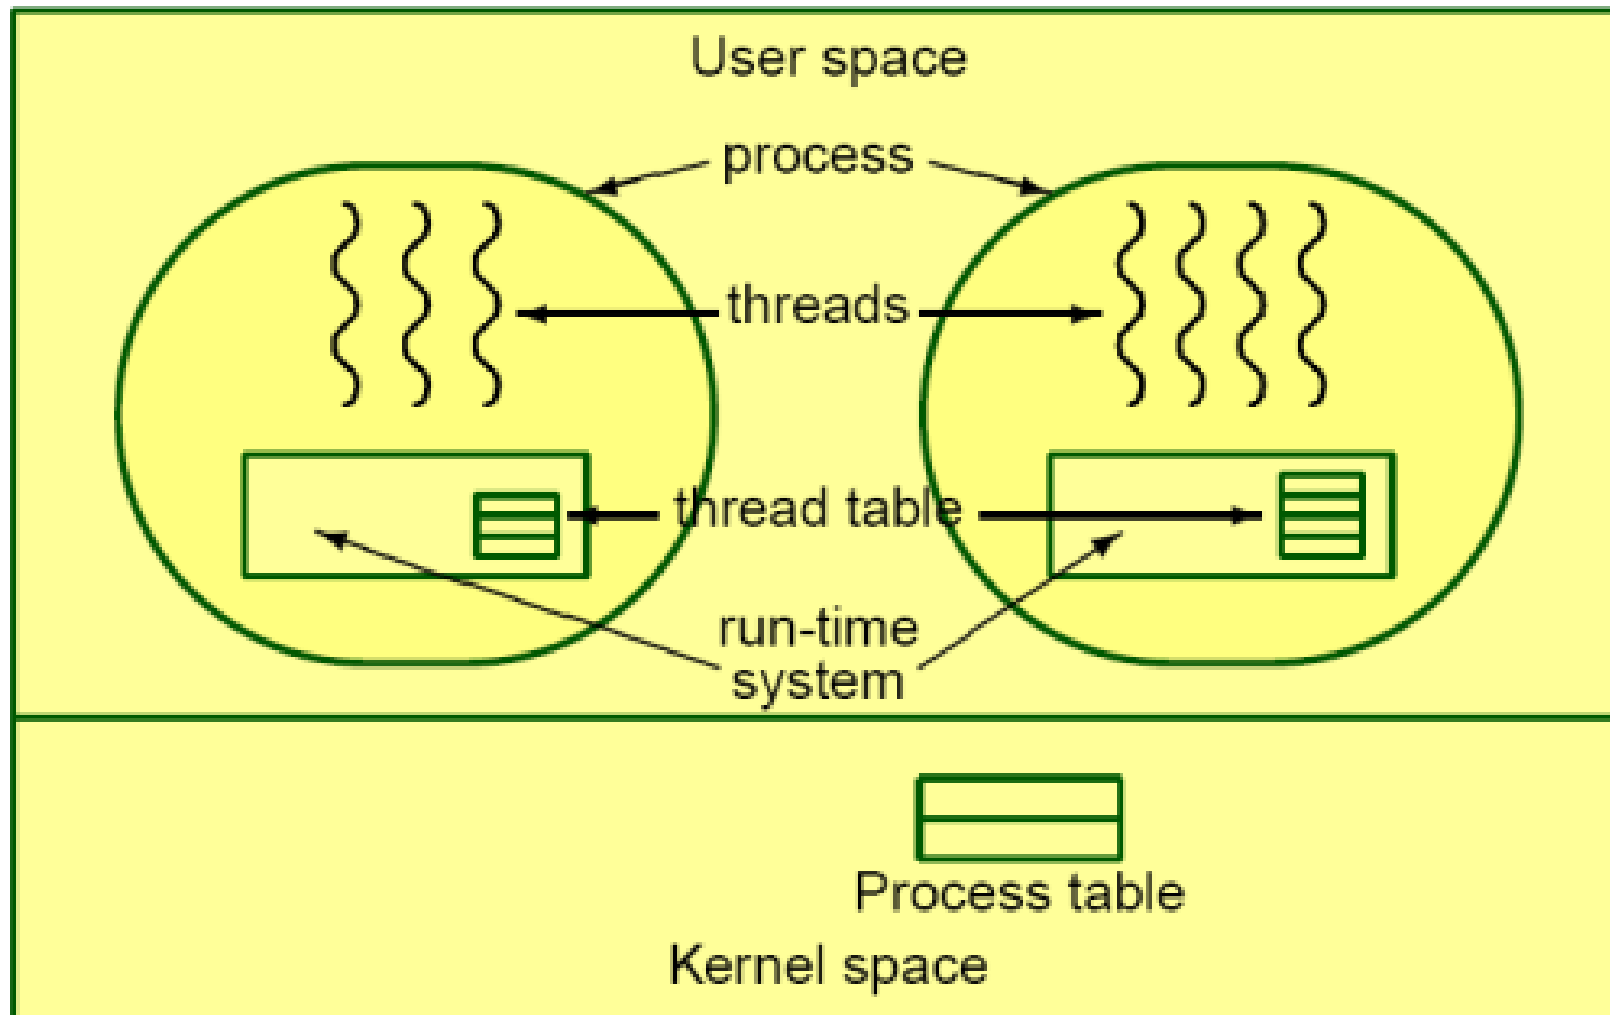
- User-level threads
  - The kernel is not aware of threads
  - Each process manages its own threads
- Kernel-level threads
  - Managed by the kernel
- Trade-offs on each side
- Various hybrid approaches possible

## User-Level Threads

- Kernel thinks it is managing processes only
- Threads implemented by software library
- Process maintains a thread table and does thread scheduling



# User-Level Threads



## Advantages of User-Level Threads

- Better performance
  - Thread creation and termination are fast
  - Thread switching is fast
  - Thread synchronization (e.g., joining other threads) is fast
  - All these operations do not require any kernel activity
- Allows application-specific run-times
  - Each application can have its own scheduling algorithm

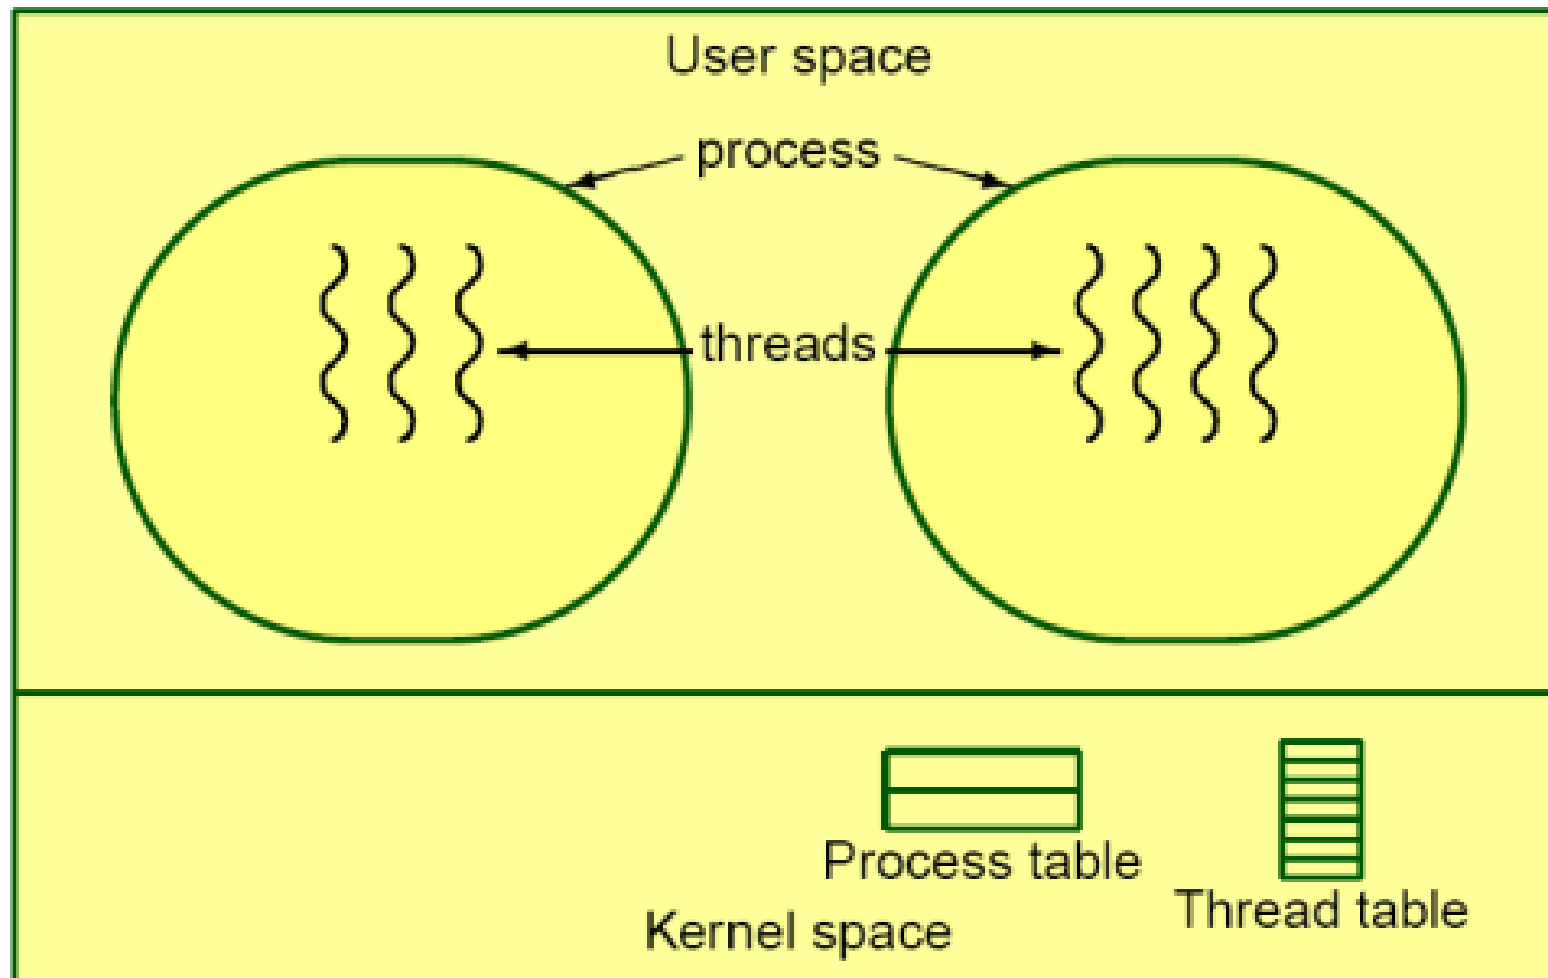
## Disadvantages of User-Level Threads

- Blocking system calls stops *all threads* in the process
  - Denies one of the core motivations for using threads
- Non-blocking I/O can be used (e.g., `select()`)
  - Harder to use and understand, inelegant
- During a page fault the OS blocks the whole process...
  - But other threads might be runnable
- Difficult to implement preemptive scheduling
  - Run-time can request a clock interrupt
    - Messy to program
    - High-frequency clock interrupts not always available
    - Individual threads may also need to use a clock interrupt

## Tutorial question

- In this problem you are to compare reading a file using a single-threaded file server and a multithreaded server, running on a single-CPU machine. It takes 15 msec to get a request for work, dispatch it, and do the rest of the necessary processing, assuming that the data needed are in the block cache. If a disk operation is needed, as is the case one-third of the time, an additional 75 msec is required, during which time the thread sleeps. For this problem, assume that thread switching time is negligible. How many requests/sec can the server handle if it is single-threaded?

# Kernel Threads



## Advantages of Kernel Threads


- Blocking system calls/page faults can be easily accommodated
  - If one thread calls a blocking system call or causes a page fault, the kernel can schedule a runnable thread from the same process

## Disadvantages of Kernel Threads

- Thread creation and termination more expensive
  - Require kernel call
  - But still much cheaper than process creation/termination
  - One mitigation strategy is to recycle threads (*thread pools*)
- Thread synchronisation more expensive
  - Requires blocking system calls
- Thread switching is more expensive
  - Requires kernel call
  - But still much cheaper than process switches
    - Same address space
- No application-specific scheduler

## Hybrid Approaches

- Use kernel threads and multiplex user-level threads onto some (or all) kernel threads





## Tutorial

- If in a multithreaded web server the only way to read from a file is the normal blocking `read()` system call, do you think user-level threads or kernel-level threads are being used? Why?