

Operating Systems

Synchronisation – Part I

Process Synchronisation


- How do processes synchronise their operation to perform a task?
- Key concepts:
 - Critical sections
 - Mutual exclusion
 - Atomic operations
 - Race conditions
 - Deadlock
 - Starvation
 - Synchronization mechanisms
 - Locks, semaphores, monitors, etc.
- Concepts relevant to both **processes** and **threads**

Shared Data Example

Account #1234: £10,000



Extract £1000
from account 1234



Extract £1000
from account 1234₃

Shared Data Example

```
void Extract(int acc_no, int sum)
{
    int B = Acc[acc_no];
    Acc[acc_no] = B - sum;
}
```

Acc[1234] 10,000

B = 10,000
Acc[1234] = 9000



Extract(1234, 1000)

B = 9,000
Acc[1234] = 8000



Extract(1234, 1000)

Shared Data Example

```
void Extract(int acc_no, int sum)
{
    int B = Acc[acc_no];
    Acc[acc_no] = B - sum;
}
```

Critical section!
Need **mutual exclusion**

Acc[1234] 10,000

B = 10,000

Acc[1234] = 9000

B = 10,000

Acc[1234] = 9000



Extract(1234, 1000)

Extract(1234, 1000)

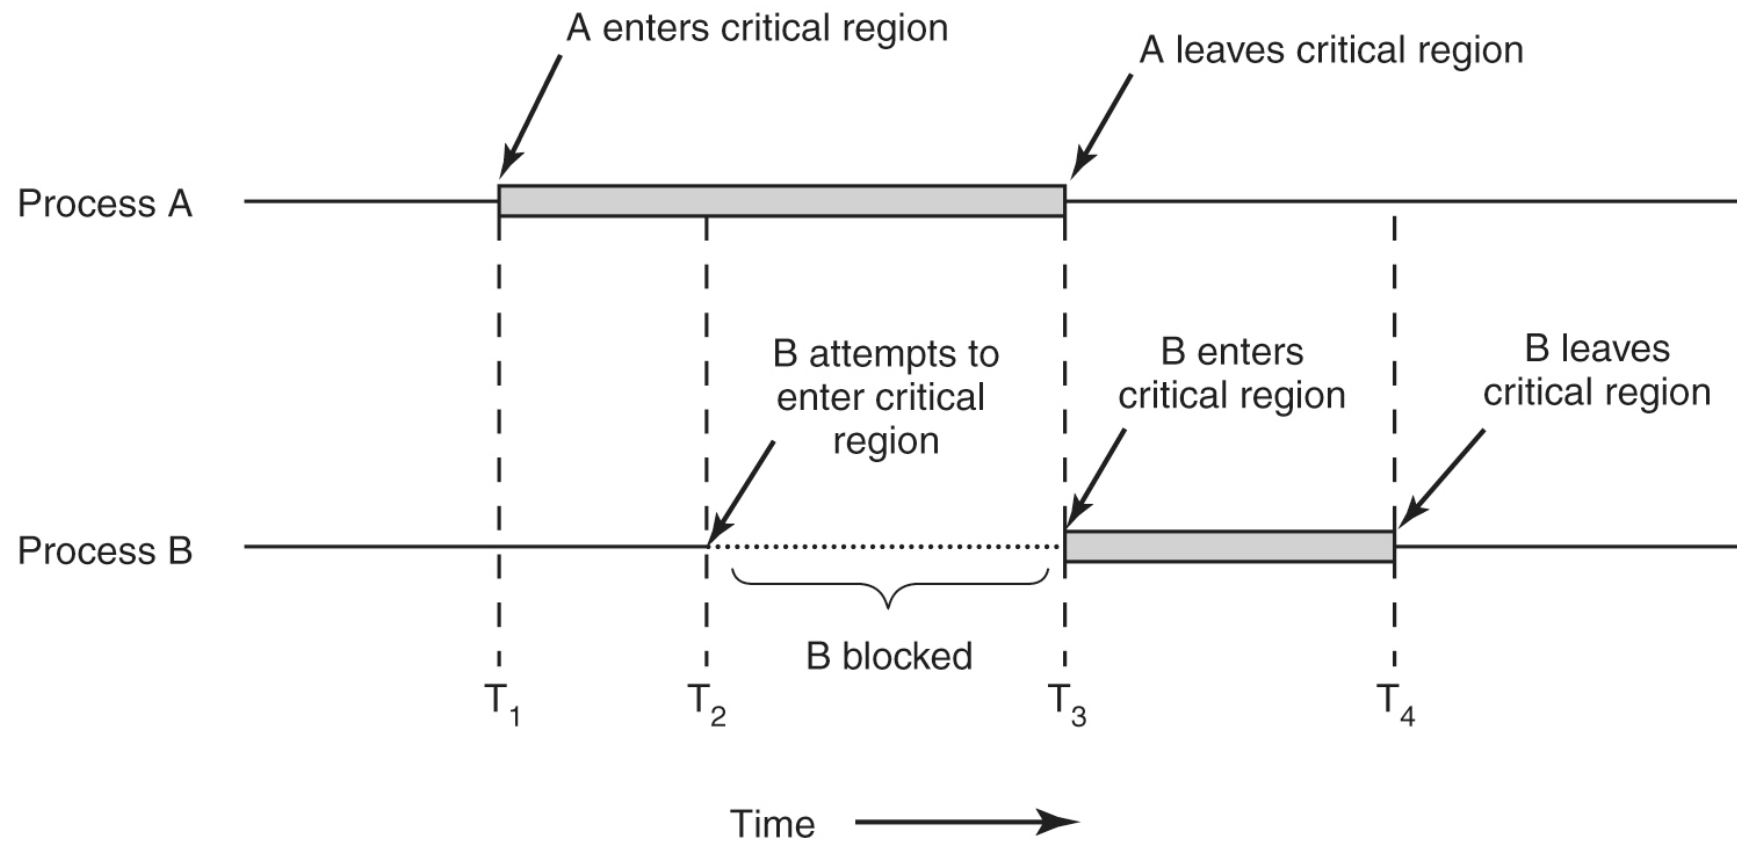
Critical Sections and Mutual Exclusion

- ***Critical section/region***: section of code in which processes access a shared resource
- ***Mutual exclusion*** ensures that if a process is executing its critical section, no other process can be executing it
 - Processes must request ***permission*** to enter critical sections
- A ***synchronisation mechanism*** is required at the entry and exit of the critical section

Requirements for Mutual Exclusion

- (1) No two processes may be simultaneously inside a critical section
- (2) No process running outside the critical section may prevent other processes from entering the critical section
 - When no process is inside a critical section, any process requesting permission to enter must be allowed to do so immediately
- (3) No process requiring access to its critical section can be delayed forever
- (4) No assumptions are made about relative the speed of processes

Critical Sections and Mutual Exclusion

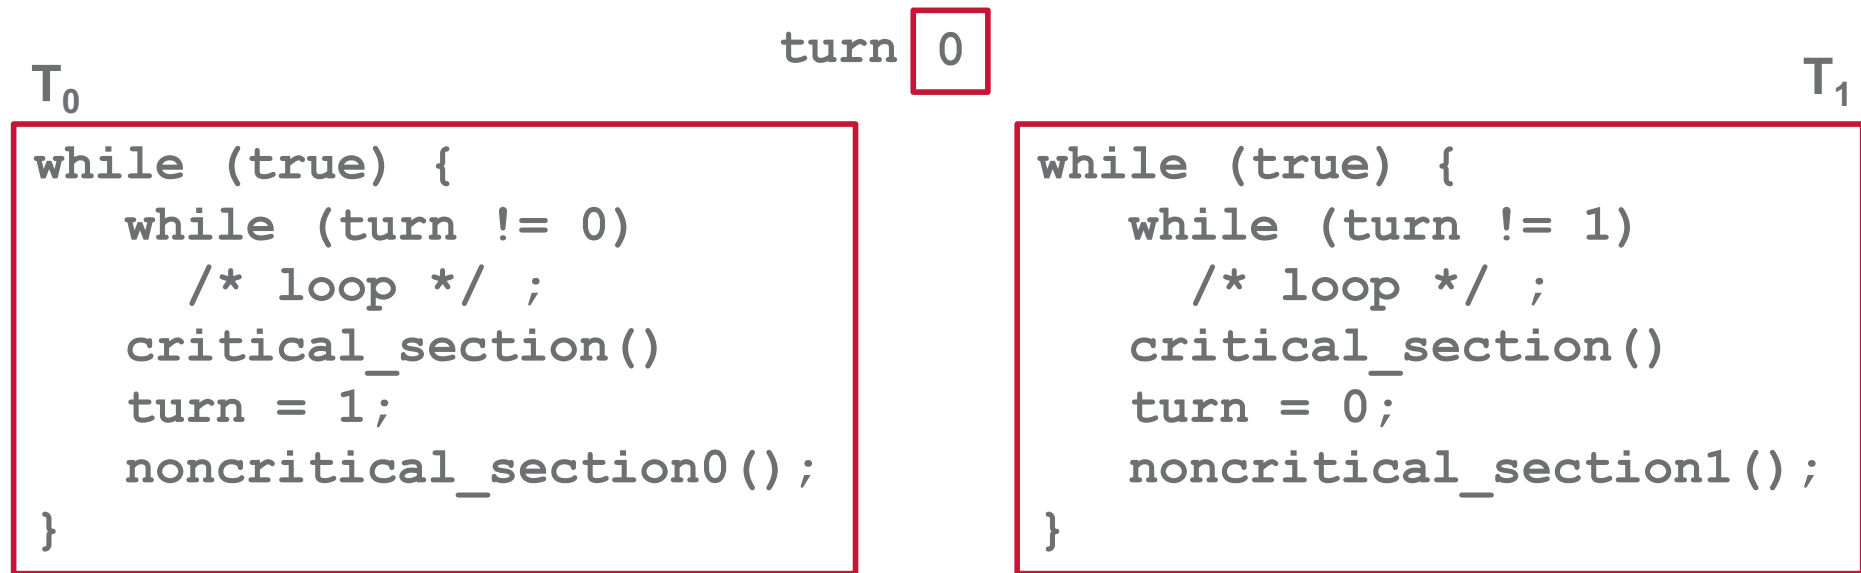


Disabling Interrupts

```
void Extract(int acc_no, int sum)
{
    CLI ();
    int B = Acc[acc_no];
    Acc[acc_no] = B - sum;
    STI ();
}
```

- Works only on single-processor systems
- Misbehaving/buggy processes may never release CPU
 - Mechanism usually only used by kernel code

Software Solution – Strict Alternation



- What happens if T_0 takes a long time in its non-critical section?
 - Remember: No process running outside its critical section may prevent other processes from entering the critical section
- Can we have T_1 execute its loop twice in a row (w/o T_0 executing in-between)?

Busy Waiting

- Strict alternation solution requires continuously testing the value of a variable
- Called **busy waiting**
 - Wastes CPU time
 - Should only be used when the wait is expected to be short

Peterson's Solution

```
int turn = 0;
int interested[2] = {0, 0};

// thread is 0 or 1
void enter_critical(int thread)
{
    int other = 1 - thread;
    interested[thread] = 1;
    turn = other;
    while (turn == other &&
           interested[other])
        /* loop */ ;
}

void leave_critical(int thread)
{
    interested[thread] = 0;
}
```

T0

```
enter_critical(0);
critical_section();
leave_critical(0);
```

T1

```
enter_critical(1);
critical_section();
leave_critical(1);
```

Peterson's Solution – Mutual Exclusion Proof

```
int turn = 0;
int interested[2] = {0, 0};

// thread is 0 or 1
void enter_critical(int thread)
{
    int other = 1 - thread;
    interested[thread] = 1;
    turn = other;
    while (turn == other &&
           interested[other])
        /* loop */ ;
}

leave_critical(int thread)
{
    interested[thread] = 0;
}
```

- First note that when T_k tries to enter or is in CS, $\text{interested}[k]=1$.
- Assume both T_0 and T_1 try to enter CS. Then:
 $\text{interested}[0]=\text{interested}[1]=1$
and turn allows only one thread to enter.
- Assume T_0 is in CS. T_1 has to wait for T_0 to set $\text{interested}[0]$ to 0, which only happens when T_0 leaves CS.
- Assume T_1 is in CS. T_0 has to wait for T_1 to set $\text{interested}[1]$ to 0, which only happens when T_1 leaves CS.

Atomic Operations

```
void Extract(int acc_no,  
            int sum)  
{  
    int B = Acc[acc_no];  
    Acc[acc_no] = B - sum;  
}
```



```
void Extract(int acc_no,  
            int sum)  
{  
    Acc[acc_no] -= sum;  
}
```

- Does this work?
 - Not atomic!
- Atomic operation: a sequence of one or more statements that is/appears to be indivisible

Lock Variables

```
void Extract(int acc_no, int sum)
{
    lock(L) ;
    int B = Acc[acc_no] ;
    Acc[acc_no] = B - sum ;
    unlock(L) ;
}
```

```
void lock(int L)
{
    while (L != 0)
        /* wait */ ;
    L = 1 ;
}
```

```
void unlock(int L)
{
    L = 0 ;
}
```

- Does this work?

TSL (Test and Set Lock) Instruction

- **Atomic** instruction provided by most CPUs
- **TSL (LOCK)**
 - Atomically sets memory location **LOCK** to 1 and returns old value

```
void lock(int L)
{
    while (TSL(L) != 0)
        /* wait */ ;
}
```

- Locks using *busy waiting* are called **spin locks**

Pseudocode; needs to be written in ASM

Spin Locks

- Waste CPU
 - Should only be used when the wait is expected to be short
- May run into **priority inversion problem**

Priority Inversion Problem and Spin Locks

- Two processes:
 - H with high priority
 - L with low priority
 - H should always be scheduled if runnable
- Assume the following scenario:
 - H is waiting for I/O
 - L acquires lock A and enters critical section
 - I/O arrives and H is scheduled
 - H tries to acquire lock A that L is holding
- What happens?

Lock Granularity

```
void Extract(int acc_no, int sum)
{
    lock(L) ;
    int B = Acc[acc_no] ;
    Acc[acc_no] = B - sum ;
    unlock(L) ;
}
```

T1: Extract(1, 40) ;

T2: Extract(2, 40) ;

- What happens if there are concurrent accesses to *different* accounts?

Lock Granularity

```
void Extract(int acc_no, int sum)
{
    lock(L[acc_no]);
    int B = Acc[acc_no];
    Acc[acc_no] = B - sum;
    unlock(L[acc_no]);
}
```

T1: Extract(1, 40);

T2: Extract(2, 40);

Lock granularity: the amount of data a lock is protecting

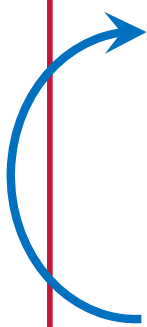
Lock Overhead and Lock Contention

- **Lock overhead:** a measure of the cost associated with using locks
 - Memory space
 - Initialization
 - *Time required to acquire and release locks*
- **Lock contention:** a measure of the number of processes waiting for a lock
 - More contention, less parallelism
- **Coarser granularity:**
 - Lock overhead?
 - Lock contention?
 - Complexity?
- **Finer granularity:**
 - Lock overhead?
 - Lock contention?
 - Complexity?

Minimizing Lock Contention/Maximizing Concurrency

- Choose finer lock granularity
 - But understand tradeoffs
- Release a lock as soon as it is not needed
 - Make critical sections small!

```
void AddAccount(int acc_no, int balance)
{
    lock(L_Acc) ;
    CreateAccount(acc_no) ;
    lock(L[acc_no]) ;
    Acc[acc_no] = balance ;
    unlock(L[acc_no]) ;
    unlock(L_Acc) ;
}
```



Read/Write Locks

```
void ViewHistory(int acc_no)
{
    print_transactions(acc_no);
}
```

T1: ViewHistory(1234) ;

T2: ViewHistory(1234) ;

T3: ViewHistory(1234) ;

- Any locks needed?

Read/Write Locks

```
void ViewHistory(int acc_no)
{
    print_transactions(acc_no);
}
```

```
void Extract(int acc_no,
             int sum)
{
    lock(L[acc_no]);
    Acc[acc_no] -= sum;
    add_debit(acc_no, sum);
    unlock(L[acc_no]);
}
```

T1: ViewHistory(1234);

T2: ViewHistory(1234);

T3: Extract(1234, 500);

- Any extra locks needed?

Read/Write Locks

```
void ViewHistory(int acc_no)
{
    lock(L[acc_no]);
    print_transactions(acc_no);
    unlock(L[acc_no]);
}
```

```
void Extract(int acc_no,
             int sum)
{
    lock(L[acc_no]);
    Acc[acc_no] -= sum;
    add_debit(acc_no, sum);
    unlock(L[acc_no]);
}
```

T1: ViewHistory(1234);

T2: ViewHistory(1234);

T3: Extract(1234, 500);

- What if later only T1 and T2 run?
 - They cannot execute concurrently anymore!

Read/Write Locks

```
void ViewHistory(int acc_no)
{
    lock_RD(L[acc_no]);
    print_transactions(acc_no);
    unlock(L[acc_no]);
}
```

```
void Extract(int acc_no,
             int sum)
{
    lock_WR(L[acc_no]);
    Acc[acc_no] -= sum;
    add_debit(acc_no, sum);
    unlock(L[acc_no]);
}
```

- Read/write locks:

- **lock_RD(L)** → acquire L in read mode
- **lock_WR(L)** → acquire L in write mode
- In write mode, the thread has exclusive access
- Multiple threads can acquire the lock in read mode!