

Operating Systems

Synchronisation – Part II

Race Condition

- Occurs when multiple threads or processes read and write **shared data** and the final result depends on the relative timing of their execution
 - i.e. on the exact process or thread **interleaving**
- E.g., the **Extract** example → final value of account 8,000 or 9,000

Thread Interleavings

```
int a, b; // shared
void T1()
{
    a = 1;
    b = 1;
}
void T2()
{
    b = 2;
    a = 2;
}
```

a = 1	a = 1	a = 1	b = 2	b = 2	b = 2
b = 1	b = 2	b = 2	a = 2	a = 1	a = 1
b = 2	b = 1	a = 2	a = 1	a = 2	b = 1
a = 2	a = 2	b = 1	b = 1	b = 1	a = 2
(2, 2)	(2, 1)	(2, 1)	(1, 1)	(2, 1)	(2, 1)

Tutorial

Consider the following three threads:

T1

a = 1;

b = 2;

T2

b = 1;

T3

a = 2;

How many possible interleavings are there?

- a) 4
- b) 8
- c) 12
- d) 16

Tutorial

If all thread interleavings are as likely to occur, what is the probability to have $a=1$ and $b=1$ after all threads complete execution?

All thread interleavings

a = 1	a = 1	a = 1	a = 1	a = 1	a = 1
b = 2	b = 2	b = 1	a = 2	b = 1	a = 2
b = 1	a = 2	b = 2	b = 2	a = 2	b = 1
a = 2	b = 1	a = 2	b = 1	b = 2	b = 2
(2, 1)	(2, 1)	(2, 2)	(2, 1)	(2, 2)	(2, 2)

b = 1	a = 2	b = 1	a = 2	b = 1	a = 2
a = 1	a = 1	a = 1	a = 1	a = 2	b = 1
b = 2	b = 2	a = 2	b = 1	a = 1	a = 1
a = 2	b = 1	b = 2	b = 2	b = 2	b = 2
(2, 2)	(1, 1)	(2, 2)	(1, 2)	(1, 2)	(1, 2)

Memory models

In this course, we assume *sequential consistency*:

- The operations of each thread appear in program order
- The operations of all threads are executed in some sequential order atomically

But other memory models (due to hardware behaviour and compiler optimisations) exist!

Sequential consistency vs weak memory models

```
1: int flag1 = 0, flag2 = 0; // shared
2: void T1()                void T2()
3: {                        {
4:   flag1 = 1;              flag2 = 1;
5:   if (flag2 == 0);        if (flag1 == 0)
6:     critical()            critical()
7: }                        }
```

- Under sequential consistency, it's impossible for both threads to read $\text{flag1} = \text{flag2} = 0$ in their if statements
- Under weak memory models, this is possible!
- Advanced reading:

<https://www.cs.princeton.edu/courses/archive/fall10/cos597C/docs/memory-models.pdf>

Happens-before relationship

- Formulated by Leslie Lamport in 1976
- Partial order relation between events (e.g., instructions) in a trace
- Denoted by $a \rightarrow b$ where a, b are events in a trace
- Consider a, b with a occurring before b in the trace:
 - If a, b are in the same thread, then $a \rightarrow b$
 - If a is `unlock(L)` and b is `lock(L)`, then $a \rightarrow b$
(can generalise for other synchronisation mechanisms)
- Irreflexive: $\forall a, a \not\rightarrow a$
- Antisymmetric: $\forall a, b: a \rightarrow b$ then $b \not\rightarrow a$
- Transitive: $\forall a, b, c: a \rightarrow b \wedge b \rightarrow c$ then $a \rightarrow c$

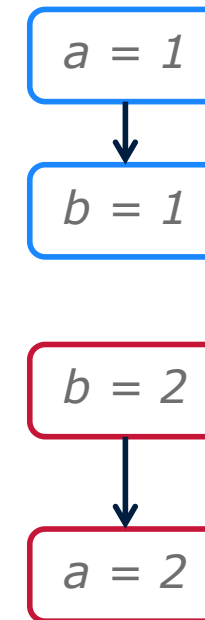
Happens-before relationship

- A data race occurs between a , b in the trace iff:
 - they access the same memory location
 - at least one of them is a write
 - they are unordered according to happens-before

Happens-before

```
int a, b; // shared
void T1()
{
    a = 1;
    b = 1;
}
```

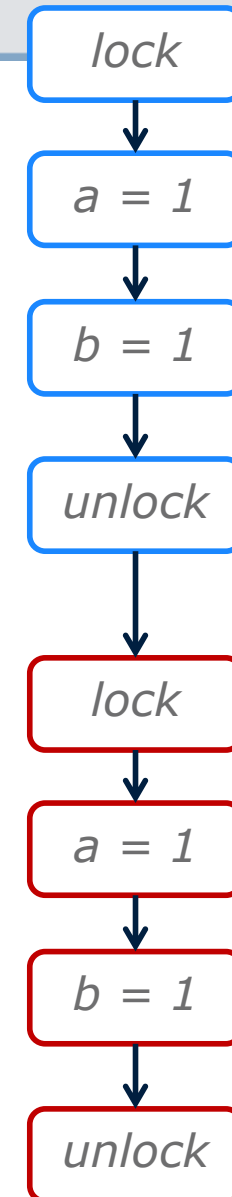
```
void T2()
{
    b = 2;
    a = 2;
}
```



Data race between $a = 1$, $a = 2$
and between $b = 1$, $b = 2$

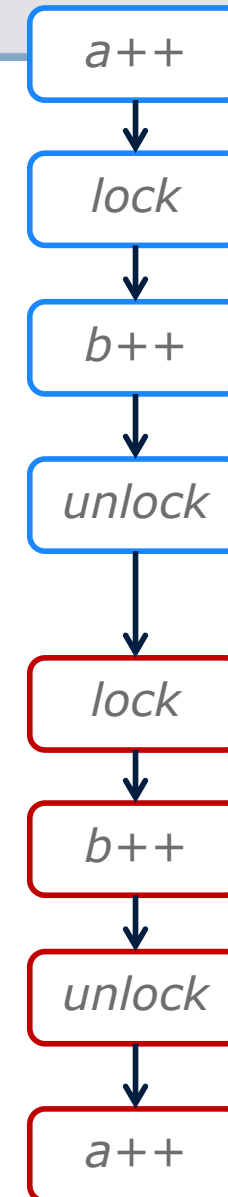
Happens-before

```
int a, b; // shared
void T1()
{
    lock(L);
    a = 1;
    b = 1;
    unlock(L);
}
void T2()
{
    lock(L);
    b = 2;
    a = 2;
    unlock(L);
}
```



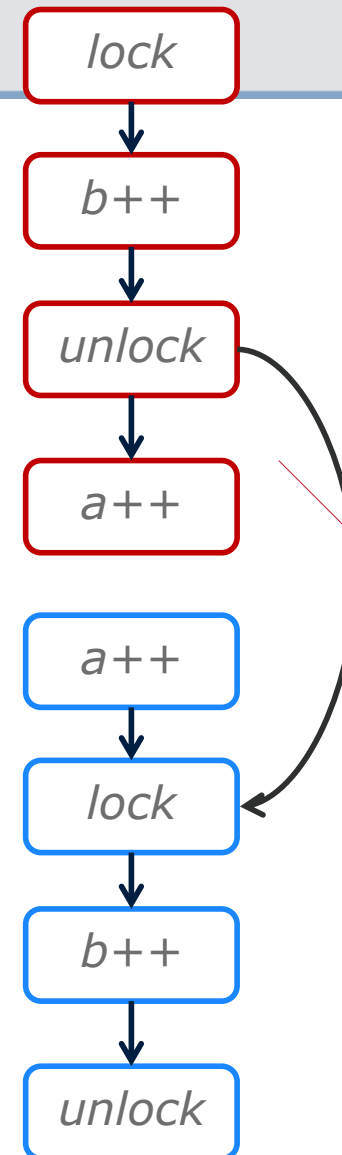
Happens-before

```
int a, b; // shared
void T1()
{
    a++;
    lock(L);
    b++;
    unlock(L);
}
void T2()
{
    lock(L);
    b++;
    unlock(L);
    a++;
}
```



Happens-before

```
int a, b; // shared
void T1()      void T2()
{              {
    a++;        lock(L)
    lock(L);    b++;
    b++;        unlock(L);
    unlock(L);  a++;
}              }
```



Semaphores

- Blocking synchronization mechanism invented by Dijkstra in 1965
- Idea: Processes will cooperate by means of *signals*
 - A process will stop, waiting for a specific signal
 - A process will continue if it has received a specific signal
- **Semaphores** are *special variables*, accessible via the following *atomic* operations:
 - **down(s)** : receive a signal via semaphore **s**
 - **up(s)** : transmit a signal via semaphore **s**
 - **init(s, i)** : initialise semaphore **s** with value **i**
- **down()** also called **P()** (*probeer te verlagen*)
- **up()** also called **V()** (*verhogen*)

Semaphores

- Semaphores have two private components:
 - A counter (non-negative integer)
 - A queue of processes currently waiting for that semaphore

Semaphore Operations

```
init(s, i) ::= counter(s) = i  
             queue(s) = {}
```

```
down(s) ::= if counter(s) > 0  
             counter(s) = counter(s) - 1  
           else  
             add P to queue(s)  
             suspend current process P
```

```
up(s) ::= if queue(s) not empty  
           resume one process in queue(s)  
         else  
           counter(s) = counter(s) + 1
```

Semaphores: Mutual Exclusion

- **Binary semaphore:** counter is initialized to 1
- Similar to a lock/mutex

```
process A                                process B
...                                     ...
down(s)                                down(s)
    critical section                    critical section
up(s)                                  up(s)
end                                    end

main() {
    var s:Semaphore
    ...
    init(s, 1) /* initialise semaphore */
    ...
    start processes A and B in random order
    ...
}
```

Semaphores: Ordering Events

Process A must execute its critical section before process B can execute its critical section

```
process A                                process B
    ...
    critical section
    up(s)
end

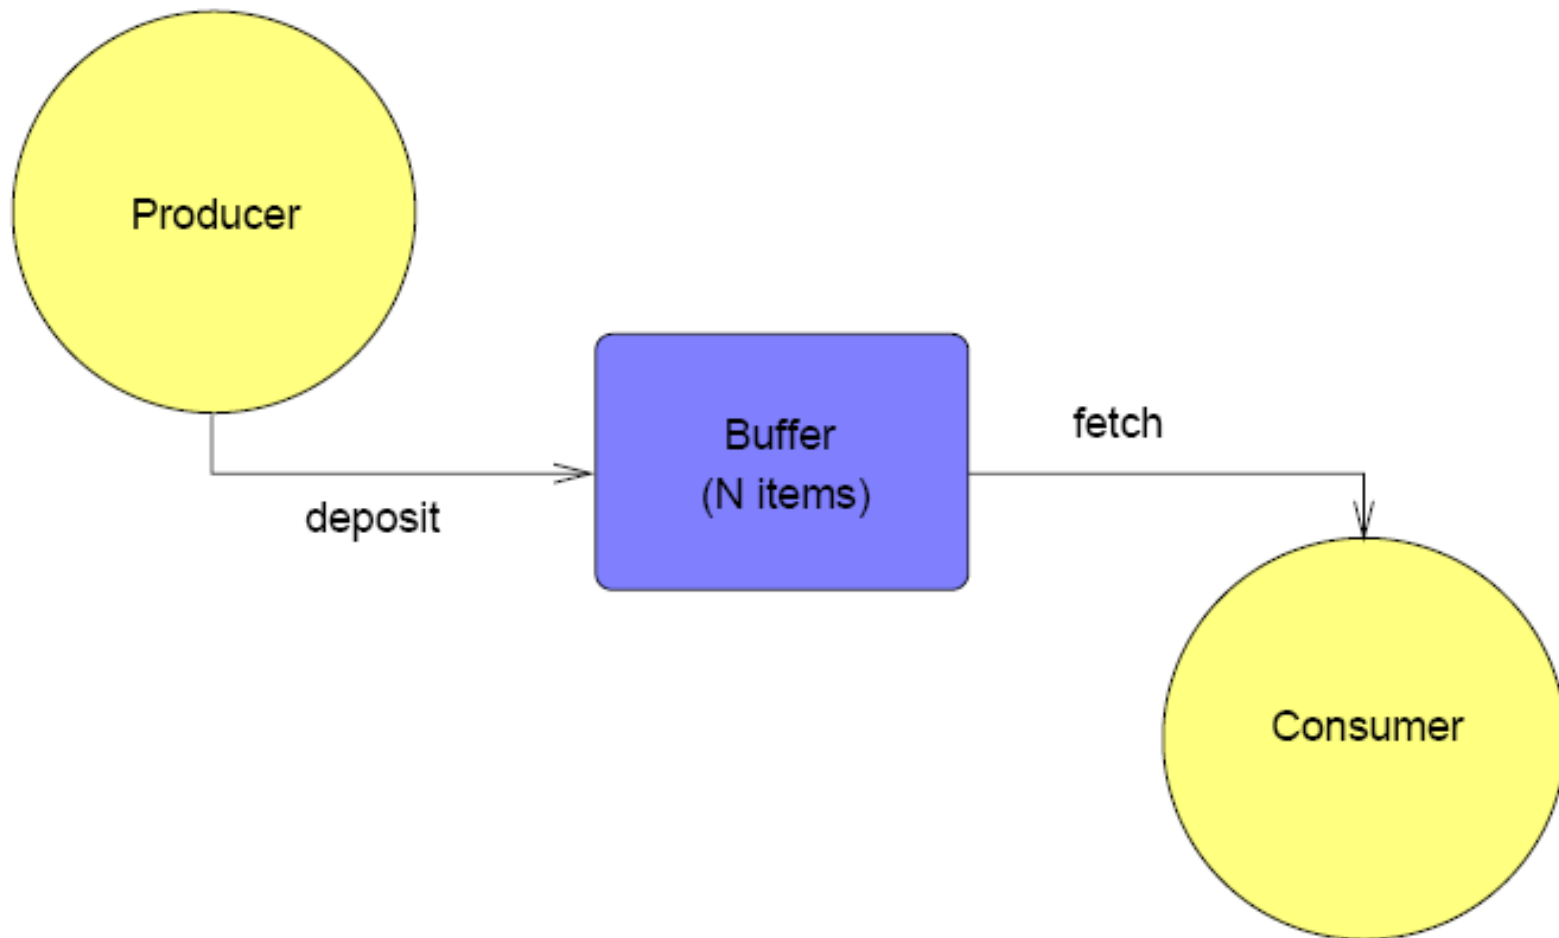
                                down(s)
                                critical section
                                end

var s:Semaphore
...
init(s, 0) /* initialise semaphore */
...
    start processes A and B in random order
...
```

General Semaphores

- The initial value of a semaphore counter indicates how many processes can access shared data at the same time
- **counter(s) ≥ 0** : how many processes can execute down without being blocked

Producer / Consumer



There can be multiple producers and consumers

Producer / Consumer

- Producer constraints:
 - Items can only be deposited in buffer if there is space
 - Items can only be deposited in buffer if mutual exclusion is ensured
- Consumer constraints:
 - Items can only be fetched from buffer if it is not empty
 - Items can only be fetched from buffer if mutual exclusion is ensured
- Buffer constraints:
 - Buffer can hold between **0** and **N** items

Producer/Consumer

```
var item, space, mutex: Semaphore
init(item, 0) /* Semaphore to ensure buffer is not empty */
init(space, N) /* Semaphore to ensure buffer is not full */
init(mutex, 1) /* Semaphore to ensure mutual exclusion */
```

```
procedure producer()
  loop
    produce item
    down(space)
    down(mutex)
    deposit item
    up(mutex)
    up(item)
  end loop
end producer
```

```
procedure consumer()
  loop
    down(item)
    down(mutex)
    fetch item
    up(mutex)
    up(space)
    consume item
  end loop
end producer
```

Monitors

- Higher-level synchronization primitive
- Introduced by Hansen (1973) and Hoare (1974)
- Refined by Lampson (1980)

Monitors

- Shared data
- Entry procedures
 - Can be called from outside the monitor
- Internal procedures
 - Can be called only from monitor procedures
- An (implicit) monitor lock
- One or more condition variables
- Processes can only call entry procedures
 - cannot directly access internal data
- **Only one process can be in the monitor at one time**

Condition Variables

- Associated with high-level conditions
 - “some space has become available in the buffer”
 - “some data has arrived in the buffer”
- Operations:
 - **wait(c)**: releases monitor lock and waits for **c** to be signalled
 - **signal(c)**: wakes up one process waiting for **c**
 - **broadcast(c)**: wakes up all processes waiting for **c**
- Signals do not accumulate
 - If a condition variable is signalled with no one waiting for it, the signal is lost

What happens on signal?

[Hoare] A process waiting for signal is immediately scheduled

- + Easy to reason about
- Inefficient: the process that signals is switched out, even if it has not finished yet with the monitor
- Places extra constraints on the scheduler

[Lampson] Sending signal and waking up from a wait not atomic

- More difficult to understand, need to take extra care when waking up from a wait()
- + More efficient, no constraints on the scheduler
- + More tolerant of errors: if the condition being notified is wrong, it is simply discarded when rechecked (see next slides)

*Usually **[Lampson]** is used (including Pintos)*

Producer/Consumer with Monitors

```
monitor ProducerConsumer
    condition not_full, not_empty;
    integer count = 0;

    entry procedure insert(item)
        if (count == N) wait(not_full);
        insert_item(item); count++;
        signal(not_empty);

    entry procedure remove(item)
        if (count == 0) wait(not_empty);
        remove_item(item); count--;
        signal(not_full);
end monitor
```

Does this work?

Producer/Consumer with Monitors

```
monitor ProducerConsumer
    condition not_full, not_empty;
    integer count = 0;

    entry procedure insert(item)
        while (count == N) wait(not_full);
        insert_item(item); count++;
        signal(not_empty);

    entry procedure remove(item)
        while (count == 0) wait(not_empty);
        remove_item(item); count--;
        signal(not_full);
end monitor
```

Monitors

- Monitors are a language construct
- Not supported by C
- Pintos
 - explicit monitor lock
- Java
 - synchronized methods
 - no condition variables
 - wait() and notify()

Recap

- Lock
 - Reader/writer locks
 - Often exposed with Monitor language construct
 - Within a process
 - 1 process/thread in critical section
- Mutex
 - Like lock, but can work across processes too
- Semaphore
 - Like mutex, but can let in N processes/threads

Bohr and Heisenbugs

- Bohrbugs:
 - Deterministic, reproducible bugs
 - Behave similar to Bohr's atom model where electrons deterministically orbit the nucleus
- Heisenbugs
 - Non-deterministic, hard to reproduce bugs
 - Often caused by race conditions
 - Suffer from the observer effect (Heisenberg Uncertainty Principle): attempts to observe them (i.e., printf's) make them disappear!
- Which bug would you rather have?
 - During development/testing: _____
 - During deployment: _____

Tutorial

- Two threads in the same process can synchronize using a kernel semaphore:
 - (1) Only if they are implemented by the kernel
 - (2) Only if they are implemented in user space
 - (3) Both if implemented by the kernel or in user-space