

# Operating Systems

## Deadlocks

# Deadlocks

- Example: two processes want to scan a document, and then save it on a CD


P0

```
down(scanner);  
down(cd_writer);  
scan_and_record();  
up(cd_writer);  
up(scanner);
```

P1

```
down(cd_writer);  
down(scanner);  
scan_and_record();  
up(scanner);  
up(cd_writer);
```

# Dining Philosophers



# Dining Philosophers

```
var chopstick:  array [0..4] of Semaphore

procedure philosopher(i:int)
  loop
    down(chopstick[i])
    down(chopstick[i+1 mod 5])
    eat
    up(chopstick[i])
    up(chopstick[i+1 mod 5])
    think
  end loop
end philosopher
```

Does this work?

What if everybody takes  
chopstick[i] at same time?

# Deadlock

- Set of processes is deadlocked if each process is **waiting for an event** that only **another process** can cause
- Resource deadlock is most common, 4 conditions must hold:
  1. **Mutual exclusion:** each resource is either available or assigned to exactly one process
  2. **Hold and wait :** process can request resources while it holds other resources earlier
  3. **No preemption:** resources given to a process cannot be forcibly revoked
  4. **Circular wait:** two or more processes in a circular chain, each waiting for a resource held by the next process

## Question

- Can the set of processes deadlocked include processes that are not in the circular chain in the corresponding resource allocation graph?

## Question


- Can a single-processor system have no processes ready and no process running? Is this a deadlocked system?

## Resource Allocation Graphs

- **Directed graph** models resource allocation
  - Directed arc from resource to process means that the process is currently owning that resource
  - Directed arc from process to resource means that the process is currently blocked waiting for that resource
- **Cycle = deadlock**



## Dining Philosophers – Deadlock Cycle



# Strategies For Dealing With Deadlock

- Ignore it
  - “The Ostrich Algorithm”
  - Contention for resources is low → deadlocks infrequent
- Detection and recovery
  - After system is deadlocked (1) detect the deadlock and (2) recover from it
- Dynamic avoidance
  - Dynamically consider every request and decide whether it is safe to grant it; needs some information regarding potential resource use
- Prevention
  - Prevent deadlocks by ensuring at least one of the four deadlock conditions can never hold


## Detection and Recovery

- Detects deadlock and recovers *after the fact*
- Dynamically builds resource ownership graph and looks for cycles
- When an *arc* has been inspected it is marked and not visited again

1. For each node do:
2. Initialise **L** to the empty list
3. Add the current node to **L** and check if it appears in **L** two times. Yes: cycle!
4. From current node check if any unmarked outgoing arch  
Yes: goto 5, No: goto 6
5. Pick unmarked outgoing arc, mark it, follow it to new current node and goto 3
6. If this is initial node then no cycles detected, terminate  
else reached dead end, remove it, go back to previous node and make it current and goto 3

*We are doing a depth-first search from each node in the graph, checking for cycles.*


## Detection – Example



**A** holds **R** and wants **S**  
**B** holds nothing wants **T**  
**C** holds nothing wants **S**  
**D** holds **U** wants **S** and **T**  
**E** holds **T** and wants **V**  
**F** holds **W** and wants **S**  
**G** holds **V** and wants **U**

## Detection – Example

- Starting at **R**, initialise  $L = []$
- Add **R** to list and move to **A** (only possibility)
- Add **A** giving  $L = [R, A]$
- Go to **S** so  $L = [R, A, S]$
- S** has not outgoing arcs  $\rightarrow$  dead end, backtrack to **A**
- A** has no outgoing arcs, backtrack to **R**
- Restart at **A**, add **A** to  $L \rightarrow$  dead end
- Restart at **B**, follow outgoing arcs until **D**, now  $L = [B, T, E, V, G, U, D]$
- Make random choice:
  - S**  $\rightarrow$  dead end and backtrack to **D**
  - Pick **T** update  $L = [B, T, E, V, G, U, D, T]$
- Cycle: Deadlock found, STOP



## Recovery

- **Pre-emption:**
  - Temporarily take resource from owner and give to another
- **Rollback:**
  - Processes are periodically **checkpointed** (memory image, state)
  - On a deadlock, **roll back** to previous state
    - When is it going to resolve the deadlock?
- **Killing processes:**
  - Select random process in cycle and kill it!
    - E.g., Might work for compile jobs, but not for DB systems

# Strategies For Dealing With Deadlock

- Ignore it
- Detection and recovery
- **Dynamic avoidance**
  - System grants resources when it knows that it is safe to do so
- Prevention

## Banker's Algorithm (Dijkstra 1965)

- Setup with a single type of resource
- N customers have a maximum credit limit expressed in number of credit units
  - E.g., 1 credit unit = £1K
- Each customer may ask for its maximum credit at some point, use it, and then repay it
- Banker knows that all customers don't need max credit at the same time
  - So it reserves less than the sum of all credit limits



## Banker's Algorithm (Dijkstra 1965)

	Has	Max
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

- Four customers A, B, C and D
  - Credit unit = £1K
- Banker reserves only 10 (instead of 22) units

## Banker's Algorithm – Safe vs. Unsafe States

### SAFE

Has Max

A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

### UNSAFE

Has Max

A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

- Safe state:
  - Are there enough resources to satisfy **any** (maximum) request from some customer?
  - Assume that customer repays loan, and then check next customer closest to the limit, etc.
- A state is **safe** iff there exists a sequence of allocations that *guarantees* that all customers can be satisfied

## Banker's Algorithm – Safe vs. Unsafe States

Has	Max		Has	Max		Has	Max		Has	Max		Has	Max	
A	3	9	A	3	9	A	3	9	A	3	9	A	3	9
B	2	4	B	4	4	B	0	–	B	0	–	B	0	–
C	2	7	C	2	7	C	2	7	C	7	7	C	0	–
Free: 3			Free: 1			Free: 5			Free: 0			Free: 7		
<b>SAFE</b>														

Has	Max		Has	Max		Has	Max	
A	4	9	A	4	9	A	4	9
B	2	4	B	4	4	B	–	–
C	2	7	C	2	7	C	2	7
Free: 2			Free: 0			Free: 4		
<b>UNSAFE</b>								

A state is **safe** iff there exists a sequence of allocations that *guarantees* all customers can be satisfied

## Banker's Algorithm – Safe vs. Unsafe States

**SAFE**

	Has	Max
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

**UNSAFE**

	Has	Max
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

- Request granted only if it leads to a safe state
- Unsafe state does not have to lead to deadlock, but banker cannot rely on this behaviour
- Algorithm can be generalized to handle multiple resource types

# Strategies For Dealing With Deadlock

- Ignore it
- Detection and recovery
- Dynamic avoidance
- **Prevention**
  - *Attack one of the four deadlock conditions:*
    - *Mutual exclusion,*
    - *Hold and wait*
    - *No preemption*
    - *Circular wait*


## Deadlock : 4 conditions

1. **Mutual exclusion:** each resource is either available or assigned to exactly one process
2. **Hold and wait :** process can request resources while it holds other resources earlier
3. **No preemption:** resources given to a process cannot be forcibly revoked
4. **Circular wait:** two or more processes in a circular chain, each waiting for a resource held by the next process

# Deadlock Prevention

- Attacking the Mutual Exclusion condition
  - E.g., share the resource
- Attacking the Hold and Wait condition
  - Require all processes to request resources before start
    - If not all available then wait
  - Issue: need to know what you need in advance
- Attacking the No-Preemption condition
  - E.g., forcing a process to give up printer half way through not good
- Attacking Circular Wait condition
  - Force single resource per process?
    - Optimality issues
  - Number resources, processes must ask for resources in this order
    - Issue: large number of resources...can be difficult to organise

## Dining Philosophers – Ordering Resources



The process holding the highest resource will never ask for a resource already assigned.



## Communication Deadlock

- E.g., process **A** sends message to **B** and blocks waiting on **B**'s reply
- **B** didn't get **A**'s message then **A** is blocked and **B** is blocked waiting on message → **deadlock!**
- Ordering resources, careful scheduling not useful here
- What should we use?
  - Communication protocol based on timeouts

# Livelock

- **Livelock**: Processes/threads are not blocked, but they or the system as a whole not making progress
- Example 1: **Enter\_region()** tests mutex then either grabs resource or reports failure. If attempt fails, it tries again.

```
process_A
{
    Enter_region (resource1)
    Enter_region (resource2)
    Use(resource1, resource2)
    Leave_region (resource2)
    Leave_region (resource1)
}
```

```
process_B
{
    Enter_region (resource2)
    Enter_region (resource1)
    Use(resource1, resource2)
    Leave_region (resource1)
    Leave_region (resource2)
}
```

- Example 2: System receiving and processing incoming messages. processing thread has lower priority and never gets a chance to run under high load (**receive livelock**)

# Starvation

- Concerns policy
- Who gets what resource when
- Many jobs want printer, who gets it?
  - Smallest file? Suits majority, fast turnaround, but what about occasional large job?
  - FCFS is more fair in this case

## Deadlocks: summary

- Resource deadlock: set of processes each waiting for a resource another one holds
  - Four conditions required for deadlock
  - Resource allocation graphs can be used to reason about deadlocks
- Strategies for avoiding deadlock
  - Ignoring
  - Detection and recovery
  - Dynamic avoidance
  - Prevention
- Communication deadlock: concerns policy
- Livelock: overall system makes no progress

## Question

- Two processes, A and B, each need three records, 1, 2, and 3, in a database. If A asks for them in the order 1, 2, 3, and B asks for them in the same order, deadlock is not possible. However, if B asks for them in the order 3, 2, 1, then deadlock is possible. With three resources, there are  $3! = 6$  possible combinations each process can request resources. What fraction of all combinations is guaranteed to be deadlock free?