


Operating Systems

Disk Management



Disks, SSDs, RAID, Caching

Disks have come a long way...



IBM 305 RAMAC (1956)

- First commercial hard disk: 4.4MB
- Footprint: 1.5 m²
- Price: \$160,000


Toshiba 0.85" disk (2005)

- Capacity: 4GB
- Price: <\$300



From Computer Desktop Encyclopedia
Reproduced with permission.
© 2005 Toshiba Corporation

Disk Evolution

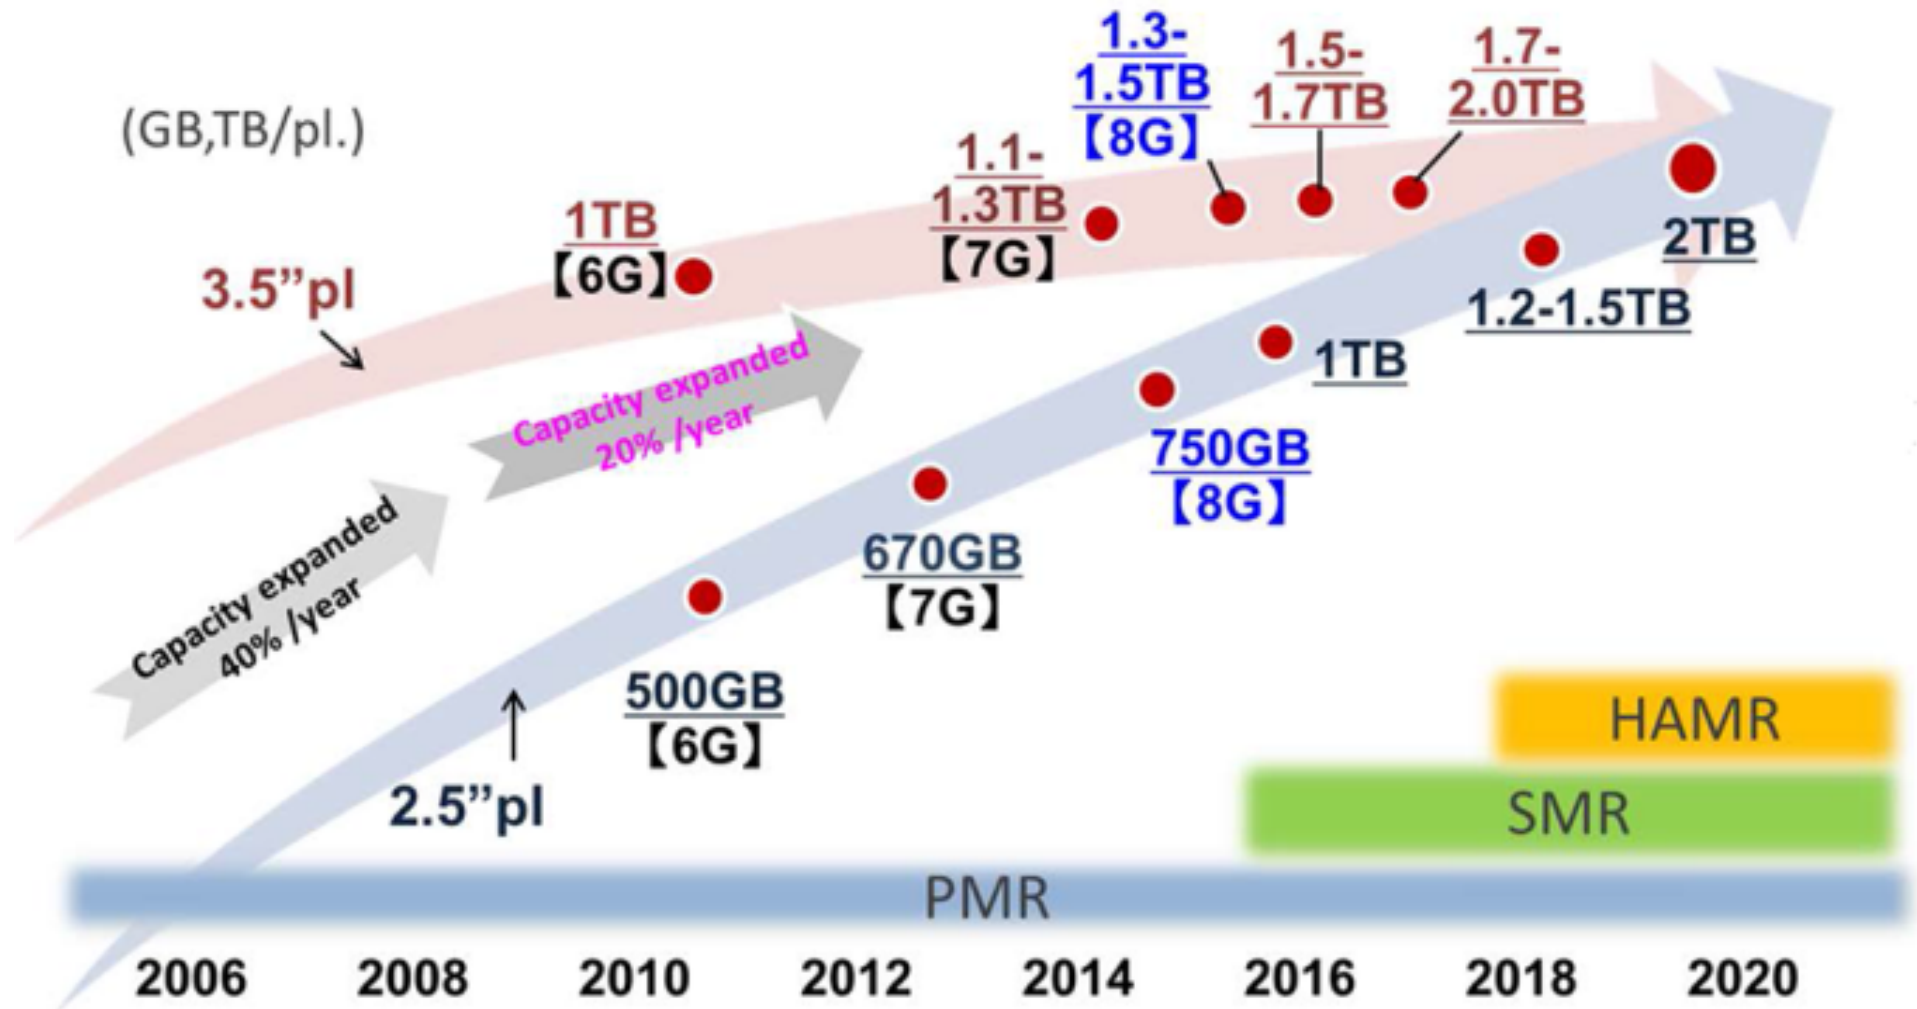


Capacity increases exponentially


- Access speeds not so much... (why?)

Disk Evolution

[Road map for storage density increase] (SDK forecast)





What is driving demand?




1GB per hour \Rightarrow 1 PB new storage *every day*
(+ multiple copies, multiple formats)

Google

Disk Storage Devices



Tracks and Cylinders




Sample Disk Specification




Parameter	IBM 360KB floppy disk	Seagate Barracuda ST3400832AS
No. of cylinders	40	16,383
Tracks / cylinder	2	16
Sectors / track	9	63
Bytes / sector	512	512
Sectors / disk	720	781,422,768
Disk capacity	360KB	400GB

Sector Layout



(a) Physical geometry



(b) Virtual geometry

Surface divided into 20 or more **zones**

- Outer zones have more sectors per track
- Ensures that sectors have same physical length
- Zones hidden using virtual geometry

Disk Addressing

Physical hardware address: (cylinder, surface, sector)

- But actual geometry complicated → hide from OS

Modern disks use **logical sector addressing** (or logical block addresses LBA)

- Sectors numbered consecutively from 0..n
- Makes disk management much easier
- Helps work around BIOS limitations
 - Original IBM PC BIOS 8GB max
 - 6 bits for sector, 4 bits for head, 14 bits for cylinder

Disk Capacity

Disk capacity statements can be confusing!

1 KB = 2^{10} bytes = 1024 bytes **vs** 1 KB = 10^3 bytes = 1000 bytes
1 MB = 2^{20} bytes = 1024² bytes **vs** 1 MB = 10^6 bytes = 1000² bytes
1 GB = 2^{30} bytes = 1024³ bytes **vs** 1 GB = 10^9 bytes = 1000³ bytes

- For the exam: just make it consistent

Disk Formatting

Before disk can be used, it must be formatted:

- **Low level format**

- Disk sector layout




- Cylinder skew
- Interleaving

- **High level format**

- Boot block
- Free block list
- Root directory
- Empty file system

Disk Delays I



Disk Delays II

Typical disk has:

<u>Sector size</u> :	512 bytes
<u>Seek time</u> (adjacent cylinder):	<1 ms
<u>Seek time</u> (average):	8 ms
<u>Rotation time</u> (average latency):	4 ms
<u>Transfer rate</u> :	up to 100MB per sec

Disk scheduling

- Minimise seek and/or latency times
- Order pending disk requests with respect to head position

Seek time approx. 2-3 larger than latency time

- More important to optimise

Disk Performance

Seek time: t_{seek}

Latency time (rotational delay): $t_{latency} = \frac{1}{2r}$

Transfer time: $t_{transfer} = \frac{b}{rN}$

where

b - number of bytes to be transferred

N - number of bytes per track

r - rotation speed in revolutions per second

Total access time: $t_{access} = t_{seek} + \frac{1}{2r} + \frac{b}{rN}$

Example: Disk Performance

Example:

Average seek time: 10ms

512 byte sectors

File size: 2560 sectors (1.3 MB)

Rotation speed: 10,000 rpm

320 sectors per track

Case A:

- Read file stored as compactly as possible on disk
 - i.e. file occupies all sectors on 8 adjacent tracks
 - $8 \text{ tracks} \times 320 \text{ sectors/track} = 2560 \text{ sectors}$

Case B:

- Read file with all sectors randomly distributed across disk

Answer: Disk Performance

Case A:

- Time to read first track

$$\text{Average seek} = 10 \text{ ms}$$

$$\text{Rotational delay} = 3 \text{ ms} = 1 / [2 * (10,000 / 60)]$$

$$\text{Read 320 sectors} = \frac{6 \text{ ms}}{19 \text{ ms}} = b / (N * (10,000 / 60))$$

- Time to read next track = $3 \text{ ms} + 6 \text{ ms} = 9 \text{ ms}$
- Total time = $19 \text{ ms} + 7 \times 9 \text{ ms} = 82 \text{ ms} = \underline{0.082 \text{ seconds}}$

Case B:

$$\text{Average seek} = 10 \text{ ms}$$

$$\text{Rotational delay} = 3 \text{ ms}$$

$$\text{Read 1 sector} = \frac{0.01875 \text{ ms}}{13.01875 \text{ ms}} = 512 / [512 * 320 * (10,000 / 60)]$$

- Total time = $2560 \times 13.01875 \text{ ms} = \underline{33.328 \text{ seconds}}$

Disk Scheduling


First Come First Served (FCFS)

No ordering of requests: random seek patterns

- OK for lightly-loaded disks
- But poor performance for heavy loads
- Fair scheduling

Queue: 98, 183, 37, 122, 14, 130, 60, 67

- Head starts at 53




Shortest Seek Time First (SSTF)

Order requests according to shortest seek distance from current head position

- Discriminates against innermost/outmost tracks
- Unpredictable and unfair performance

Queue: 98, 183, 37, 122, 14, 130, 60, 67

- Head starts at 53
- If, when handling request at 14, new requests arrive for 50, 70, 100, → long delay before 183 serviced




SCAN Scheduling

Choose requests which result in shortest seek time
in preferred direction

- Only change direction when reaching outermost/innermost cylinder (or no further requests in preferred direction)
- Most common scheduling algorithm
- Long delays for requests at extreme locations

Queue: 98, 183, 37, 122, 14, 130, 60, 67

- Head starts at 53; direction: towards 0



Sometimes called
elevator scheduling


C-SCAN

Services requests in one direction only

- When head reaches innermost request, jump to outermost request
- Lower variance of requests on extreme tracks
- May delay requests indefinitely (though less likely)

Queue: 98, 183, 37, 122, 14, 130, 60, 67

- Head starts at 53




N-Step SCAN

As for SCAN, but services only requests waiting when sweep began

- Requests arriving during sweep serviced during return sweep
- Doesn't delay requests indefinitely

Queue: 98, 183, 37, 122, 14, 130, 60, 67

- Head starts at 53; direction: towards 0.
- Requests 80, 140 arrive when head moving outwards



Linux: Disk Scheduling

I/O requests placed in **request list**

- One request list for each device in system
- `bio` structure: associates memory pages with requests

Block device drivers define **request** operation called by kernel

- Kernel passes ordered request list
- Driver must perform all operations in list
- Device drivers do not define read/write operations

Some devices drivers (e.g. RAID) order their own requests

- Bypass kernel for request list ordering (**why?**)

Linux: Disk Scheduling Algorithms

Default: variation of SCAN algorithm

- Kernel attempts to merge requests to adjacent blocks
- But: synchronous read requests may starve during large writes


Deadline scheduler: ensures reads performed by deadline

- Eliminates read request starvation

Anticipatory scheduler: delay after read request completes

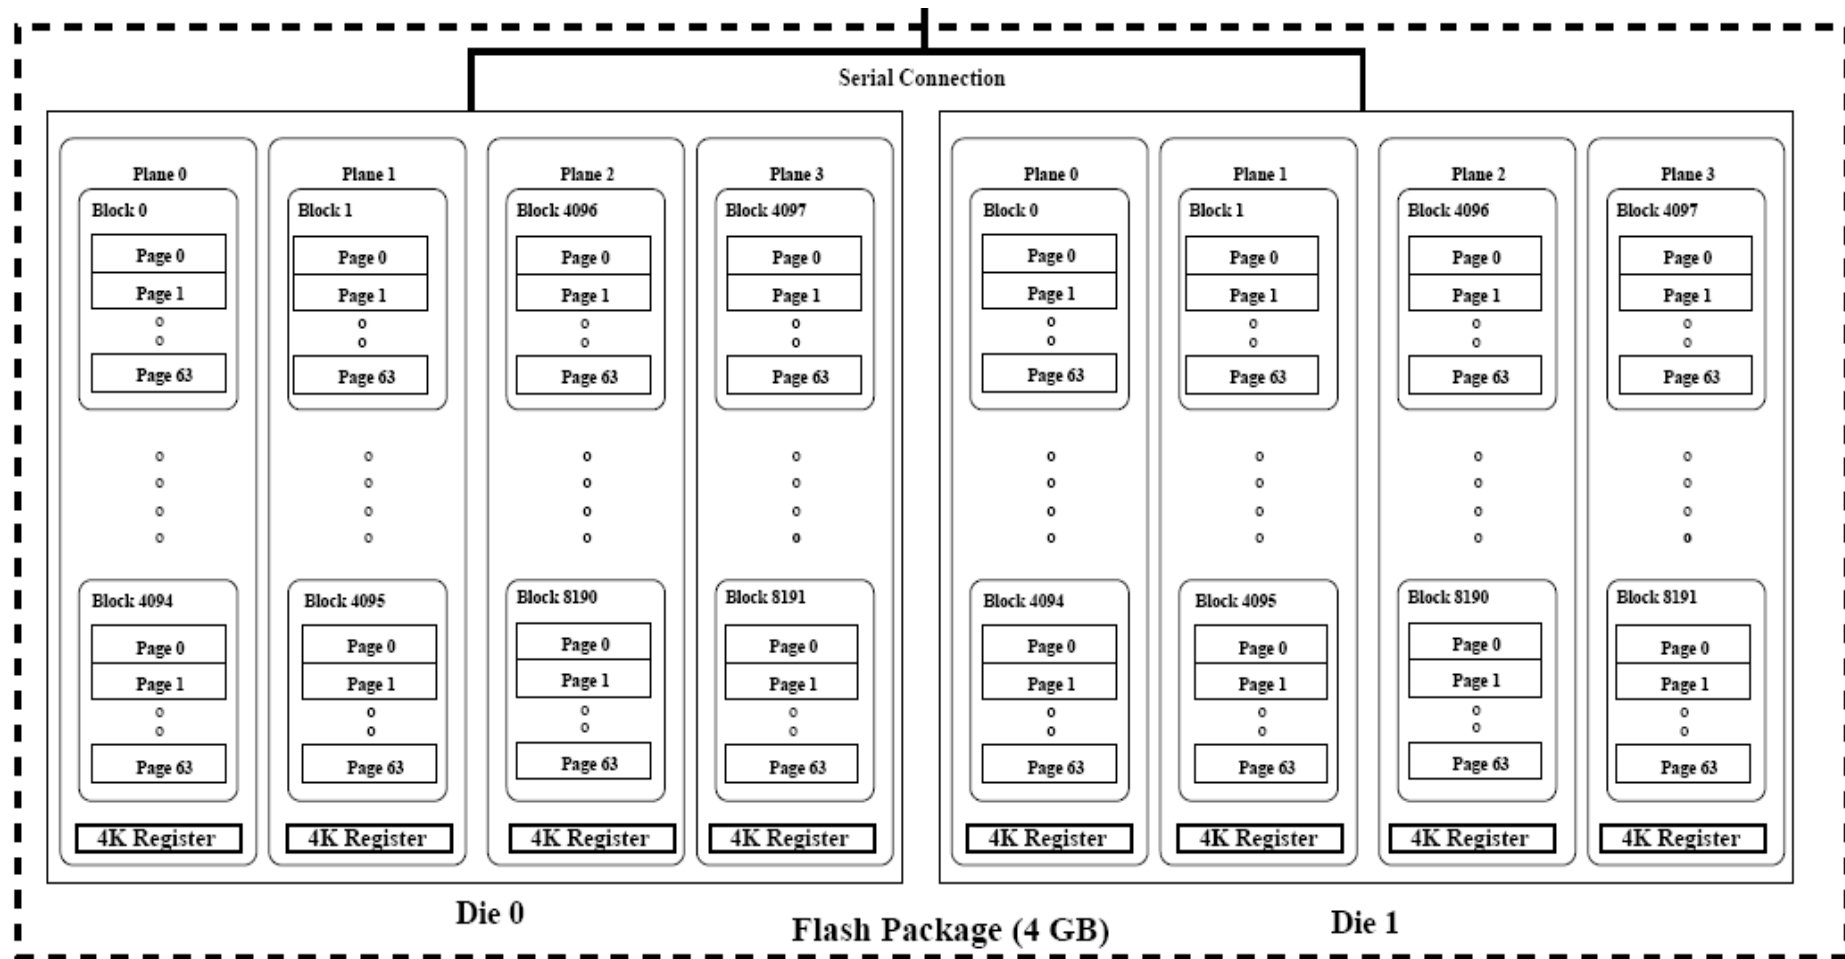
- Idea: process will issue another synchronous read operation before its quantum expires
- Reduces excessive seeking behaviour
- Can lead to reduced throughput if process does not issue another read request to nearby location
 - Anticipate process behaviour from past behaviour

Solid State Drives (SSDs)



<http://blog.digistor.com/under-the-hood-industrial-ssd-advancements/>

Internals




SSDs vs HDDs

SSDs

- More bandwidth (1GB/s read/write vs 100MB/s)
- Smaller latencies (microseconds vs milliseconds)
- So SSDs are always better?

Detailed tradeoffs



If you care about IOPS/\$, then choose SSDs
YouTube doesn't run on SSDs (2017)

RAID

RAID

Problem:

- CPU performance doubling every 18 months
- Disk performance has increased only 10 times since 1970

Solution: Use parallel disk I/O

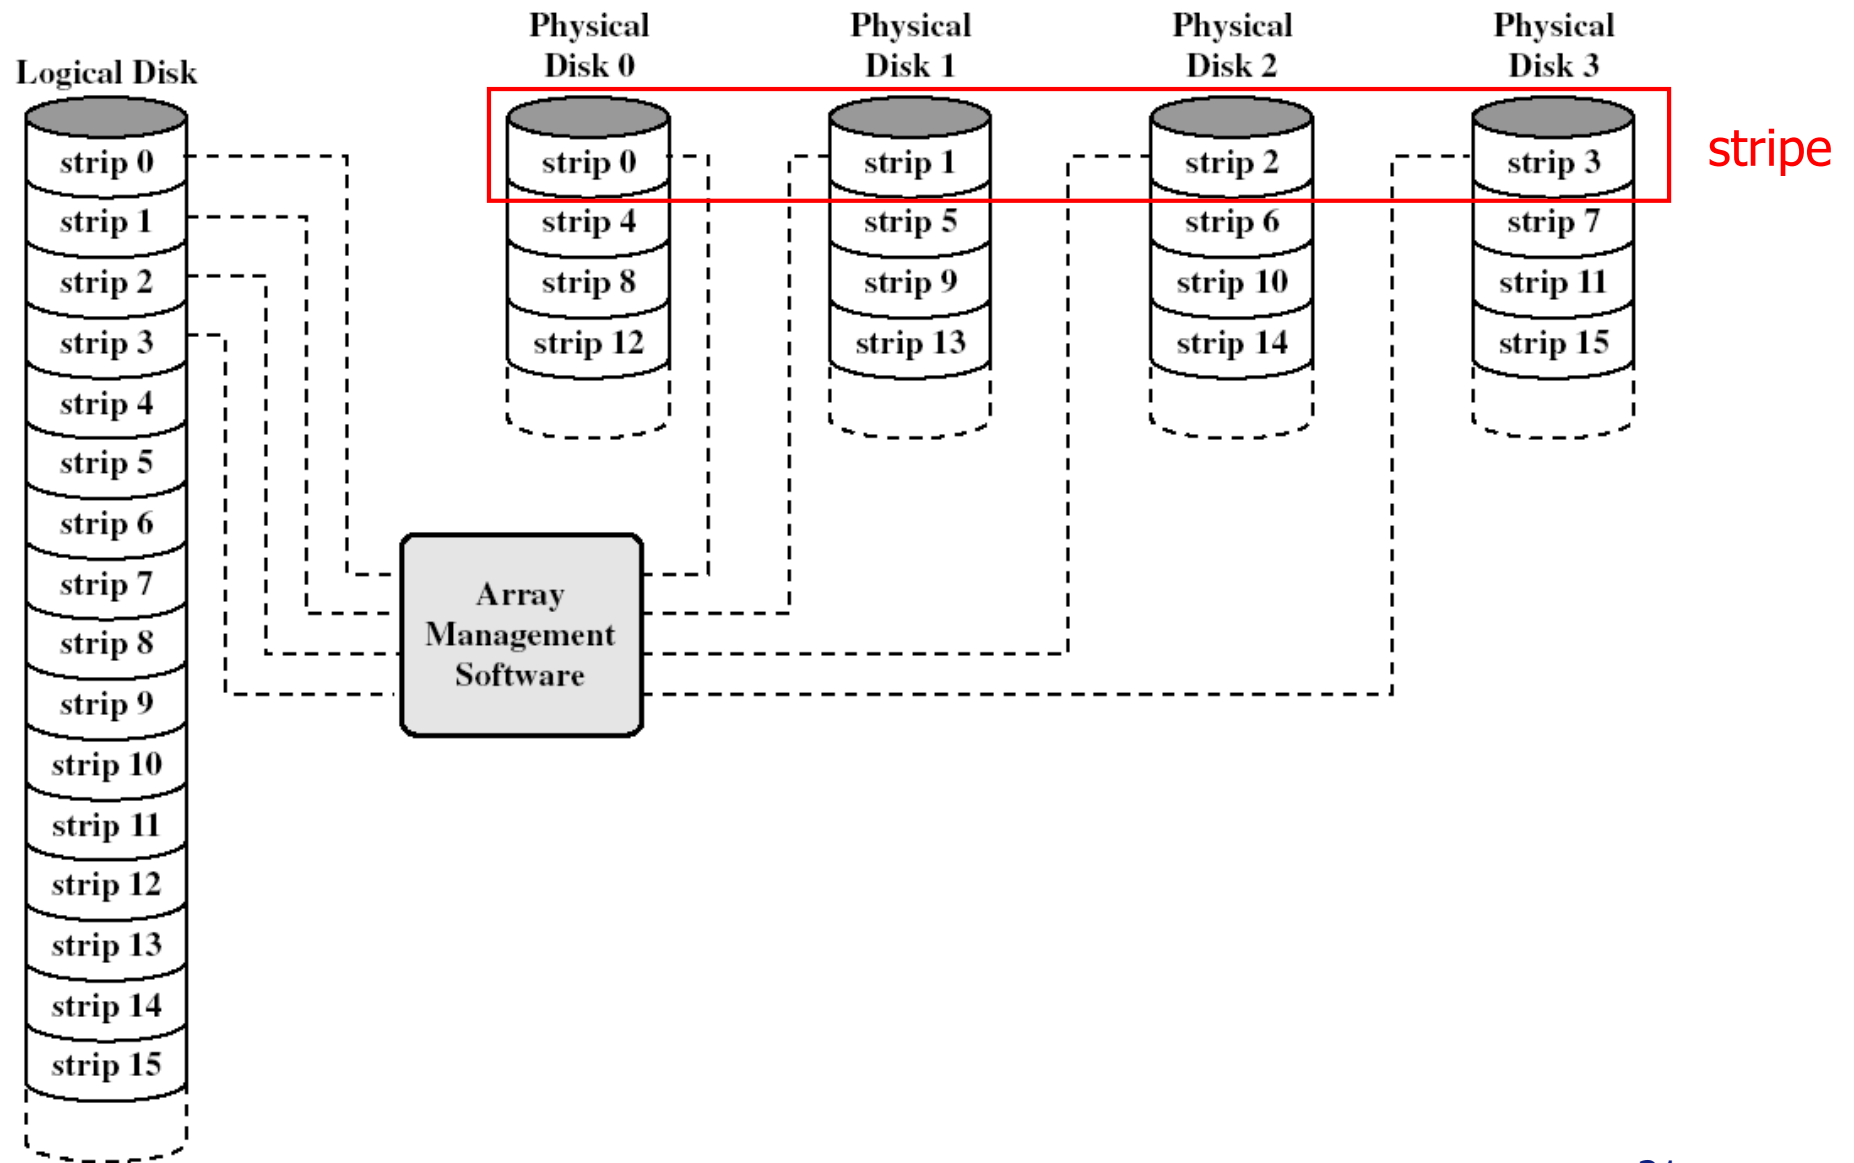
RAID (Redundant Array of Inexpensive Disks)

- Use array of physical drives appearing as single virtual drive
- Stores data distributed over array of physical disks to allow parallel operation (called **striping**)

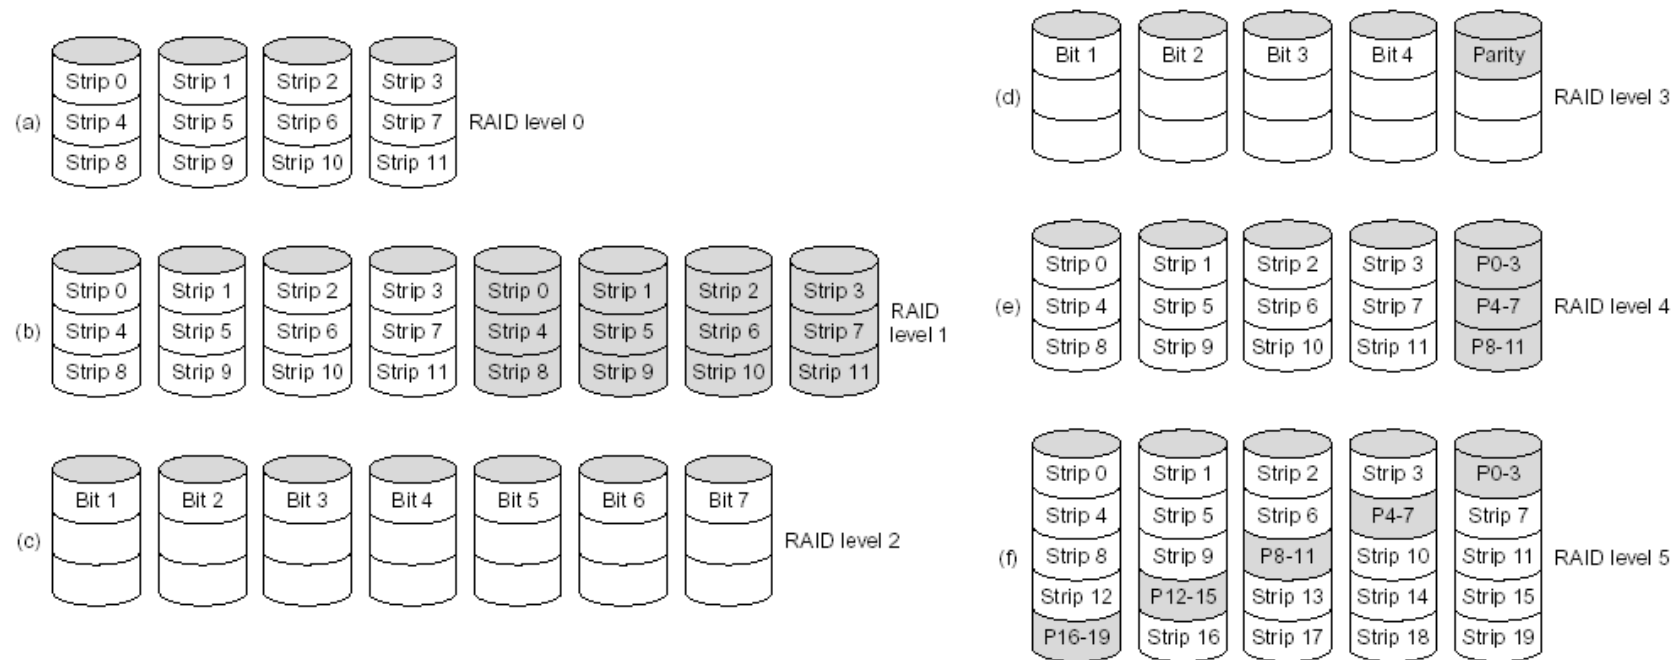
Use redundant disk capacity to respond to disk failure

- More disks \Rightarrow lower mean-time-to-failure (**MTTF**)

RAID: Striping



RAID Levels



RAID levels with different properties in terms of

- performance characteristics
- level of redundancy
- degree of space efficiency (cost)

Some other these are of historic interest...


RAID Level 0 (Striping)

Use multiple disks and spread out data

Disks can seek/transfer data concurrently

- Also may balance load across disks

No redundancy → no fault tolerance



RAID Level 1 (Mirroring)


Mirror data across disks

Reads can be serviced by either disk (fast)

Writes update both disks in parallel (slower)

Failure recovery easy

- High storage overhead (high cost)



RAID Level 2 (Bit-Level Hamming)

Parallel access by striping at bit-level


- Use Hamming error-correcting code (ECC)
- Corrects single-bit errors (and detect double-bit errors)

Very high throughput for reads/writes

- But all disks participate in I/O requests (no concurrency)
- **Read-modify-write cycle**

Only used if high error rates expected

- ECC disks become bottleneck
- High storage overhead



RAID Level 3 (Byte-Level XOR)


Only single parity strip used

Parity = data1 XOR data2 XOR data3 ...

- Reconstruct missing data from parity and remaining data

Lower storage overhead than RAID L2

- But still only one I/O request can take place at a time



RAID Level 4 (Block-Level XOR)


Parity strip handled on block basis

- Each disk operates independently

Potential to service multiple reads concurrently

Parity disk tends to become bottleneck

- Data and parity strips must be updated on each write



RAID Level 5 (Block-Level Distributed XOR)


Like RAID 4, but distribute parity

- Most commonly used

Some potential for write concurrency

Good storage efficiency/redundancy trade-off

- Reconstruction of failed disk non-trivial (and slow)



RAID Summary

Category	Level	Description	I/O Data Transfer (read/write)	I/O Request Rate (reads/writes)
Striping	0	Non-redundant	+ / +	+ / +
Mirroring	1	Mirrored	+ / 0	+ / 0
Parallel access	2	Redundant via Hamming code	++ / ++	0 / 0
	3	Bit interleaved parity	++ / ++	0 / 0
Independent access	4	Block interleaved parity	+ / -	+ / -
	5	Block interleaved distributed parity	+ / -	+ / - or 0

better than single disk (+) / same (0) / worse (-)

Disk Caching

Disk Cache

Idea: Use main memory to improve disk access


Buffer in main memory for disk sectors

- Contains copy of some sectors from disk
- OS manages disk in terms of **blocks**
 - Multiple sectors for efficiency
 - cf. Device Management (block devices)

Buffer uses finite space

- Need **replacement policy** when buffer full

Buffer Cache



Least Recently Used (LRU)

Replace block that was in cache longest with no references

Cache consists of stack of blocks

- Most recently referenced block on top of stack
- When block referenced (or brought into cache), place on top of stack
- Remove block at bottom of stack when new block brought in

Don't move blocks around in main memory

- Use stack of pointers instead

Problem: Doesn't keep track of block "popularity"

Least Frequently Used (LFU)

Replace block that has experienced fewest references

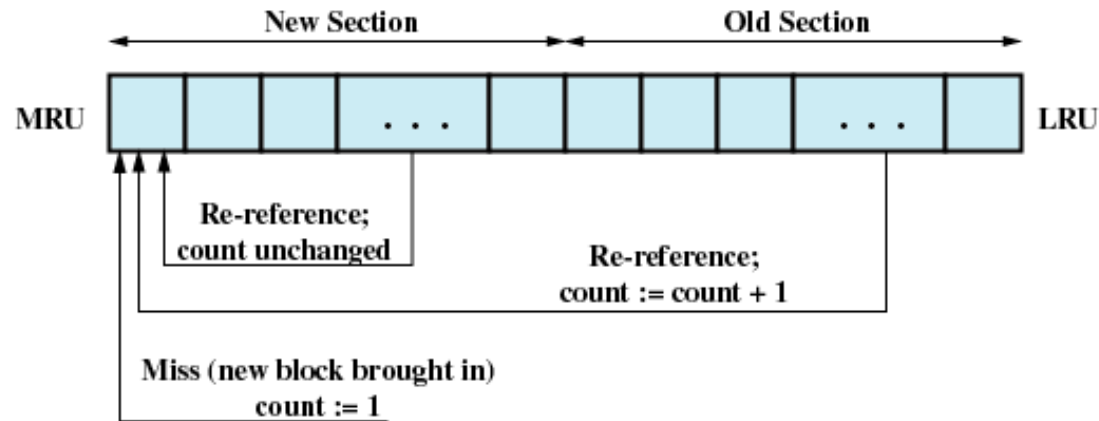
Counter associated with each block

- Counter incremented each time block accessed
- Block with smallest count selected for replacement

Some blocks may be referenced many times in short period of time

- Leads to misleading reference count
- Use **frequency-based replacement**

Frequency-Based Replacement



Divide LRU stack into two sections: new and old

- Block referenced → move to top of stack
- Only increment reference count if not already in new

Problem: blocks “age out” too quickly (why?)

- Use three sections and only replace blocks from old

